

# History of Database Systems

---

JIANNAN WANG

SIMON FRASER UNIVERSITY

# Database Systems in a Half Century

## (1960s – 2010s)

When	What
Early 1960 – Early 1970	The Navigational Database Empire
Mid 1970 – Mid 1980	The Database World War I
Mid 1980 – Early 2000	The Relational Database Empire
Mid 2000 – Now	The Database World War II

### References.

- <https://en.wikipedia.org/wiki/Database#History>
- [What Goes Around Comes Around \(Michael Stonebraker, Joe Hellerstein\)](#)
- [40 Years VLDB Panel](#)

# The Navigational Database Empire

## (Early 1960 – Early 1970)

### Data Model

1. How to organize data
2. How to access data

### Navigational Data Model

1. Organize data into a multi-dimensional space (i.e., A space of records)
2. Access data by following pointers between records

### Inventor: Charles Bachman

1. The 1973 ACM Turing Award
2. Turing Lecture: "The Programmer As Navigator"



# The Navigational Database Empire

## (Early 1960 – Early 1970)

---

### Representative Navigational Database Systems

- Integrated Data Store (IDS), 1964, GE
- Information Management System (IMS), 1966, IBM
- Integrated Database Management System (IDMS), 1973, Goodrich

### CODASYL

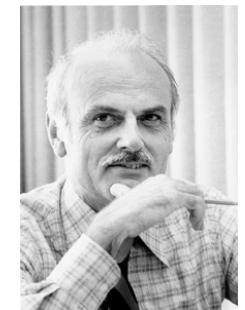
- Short for “Conference/Committee on Data Systems Languages”
- Define navigational data model as standard database interface (1969)

# The Birth of Relational Model

---

## Ted Codd

- Born in 1923
- PHD in 1965
- "A Relational Model of Data for Large Shared Data Banks" in 1970



## Relational Model

- Organize data into a collection of relations
- Access data by a declarative language (i.e., tell me what you want, not how to find it)

Data  
Independence

# The Database World War I: Background

---

## One Slide (Navigational Model)

- Led by Charles Bachman (1973 ACM Turing Award)
- Has built mature systems
- Dominated the database market

## The other Slide (Relational Model)

- Led by Ted Codd (mathematical programmer, IBM)
- A theoretical paper with no system built
- Little support from IBM

# The Database World War I: Three Big Campaigns

---

1. Which data model is better in **theory**?  
(Mid 1970)
  
2. Which data model is better in **practice**?  
(Late 1970 – Early 1980)
  
3. Which data model is better in **business**?  
(Early 1980 – Mid 1980)

# **The “Theory” Campaign**

---

A **debate** at ACM SIGFIDET (precursor of SIGMOD)  
1974

**Navigational model is bad**

- **Data Organization:** So complex
- **Data Access:** No declarative language

**Relational model is bad**

- **Data Organization:** A special case of navigational model
- **Data Access:** No system proof that declarative language is viable

# The “Practice” Campaign

---

## The Big Question

- Can a relational database system perform as good as a navigational system?

## System prototypes

- Ingres at UC Berkeley (early and pioneering) 
- System R at IBM (arguably got more stuff “right”) 

## The System R Team

- Query Optimization (Patricia P. Griffiths et al.)
- SQL ( Donald D. Chamberlin et al.)
- Transaction (Jim Gray et al.)

# The “Business” Campaign

---

Commercialization of Relational Database Systems

Not as easy as we thought

Three reasons (required) that led relational database systems to win

- The minicomputer revolution (1977)
- Competing products (e.g. IDMS) could not be ported to the minicomputer
- Relational front end was not added to navigational database systems

# What Can We Learn?

---

**Lesson 1.**

The winning of theory  $\neq$  The winning of practice

**Lesson 2.**

The winning of practice  $\neq$  The winning of business

**Lesson 3.**

Everyone can get a chance to win

# The Relational Database Empire (Mid1980 – Early 2000)

---

## Parallel and distributed DBs (1980 – 1990)

- SystemR\*, Distributed Ingres, Gamma, etc.

## Object-oriented DB (1980 – 1990)

- Objects: Data/Code Integration
- Extensibility: User-defined functions, User-defined data types

## MySQL and PostgreSQL (1990s)

- Widely used open-source relational DB systems

# The Database World War II: Background

---

## Internet Boom (Early 2000)

- Larger data volume that cannot be fit in a single machine
- Faster data updates that cannot be handled by a single machine

Commercial distributed database systems are expensive 😞

Open-source database systems do not support distributed computing well 😞

# **OLTP**

OnLine Transaction Processing

## **Workload**

High-frequent Updates + Small Queries

# **OLAP**

OnLine Analytical Processing

## **Workload**

Low-frequent Updates + Big Queries

# The Database World War II: Two Big Campaigns

---

1. Which is better for OLTP:  
NoSQL vs. Relational DBMS?
2. Which is better for OLAP:  
MapReduce vs. Relational DBMS?

# The Database World War II: Two Big Campaigns

---

- 1. Which is better for OLTP:  
NoSQL vs. Relational DBMS?**
- 2. Which is better for OLAP:  
MapReduce vs. Relational DBMS?**

# What Happened To OLTP?

---

OldSQL (1970 – Now)

NoSQL (2000 – Now)

NewSQL (2010 – Now)

# OldSQL (1970 – Now)

---

## Traditional SQL vendors



...

Still very big market!!!

Limitation 1: Not Scalable

Limitation 2: Pre-defined Schema

# The advent of Web 2.0

---

**Read-only Web → Read-write Web**



**Highly Scalable**

- Scale to 1,000,000 users and 1000 servers

**Highly Available**

- Available 24 hours a day, 7 days a week

**Highly Flexible**

- Flexible schema and flexible data types

# NoSQL Pioneers

---

**Memcached** [Fitzpatrick 2004]

- **In-memory indexes** can be highly scalable

**BigTable** [Chang et al. 2006]

- **Persistent record storage** could be scaled to thousands of nodes

**Dynamo** [DeCandia et al. 2007]

- **Eventual consistency** allows for higher availability and scalability

# NoSQL Categories

---

NoSQL	Data Model	Example Systems
Key-value Stores	Hash	DynamoDB, Riak, Redis, Membase
Document Stores	Json	SimpleDB, CouchBase, MongoDB
Wide-column Stores	Big Table	Hbase, Cassandra, HyperTable
Graph Database	Graph	Neo4J, InfoGrid, GraphBase

# NoSQL Limitations

---

## Low-level Language

- Simple read/write database operators

## Weak Consistency

- Eventual Consistency

## Lack of Standardization

- 100+ NoSQL systems

# NewSQL

---

## Strong Consistency + High Scalability



The end of an architectural era:(it's time for a complete rewrite)

M Stonebraker, S Madden, DJ Abadi... - Proceedings of the 33rd ..., 2007 - dl.acm.org

Abstract In previous papers [SC05, SBC+ 07], some of us predicted the end of "one size fits all" as a commercial relational DBMS paradigm. These papers presented reasons and experimental evidence that showed that the major RDBMS vendors can be outperformed ...

Cited by 580 Related articles All 55 versions Cite Save

90% query time spent on overhead

# NewSQL Market

---

**ScaleBase**  
Scaling Your Data At The Speed Of Your Business

**MySQL**™ Cluster

**HEKATON**  
Microsoft®  
SQL Server

**SAP** HANA

**Clustrix**

**memsql**

**VOLTDB**

**Pivotal**™

## Limitations

- Scalable but not highly scalable
- Available but not highly available
- Flexible but not highly flexible

# The Database World War II: Two Big Campaigns

---

- 1. Which is better for OLTP:  
NoSQL vs. Relational DBMS?**
- 2. Which is better for OLAP:  
MapReduce vs. Relational DBMS?**



Secure

<https://www.google.com/about/our-company/>

“Organize the world’s information and make it universally accessible and useful.”

**1 PB = 100GB \*10,000 Machines**

How to store 1PB using 10,000 machines?

GFS (HDFS)

How to process 1PB using 10,000 machines?

MapReduce

# Why MapReduce?

## 1. Fault Tolerant

(Your program will be OK when failures happen)

# of machines	Failure Probability
1	0.1%
10	0.9%
100	9.5%
1000	63.2%
10,000	??? 99.9%

Cost for 5 nodes	Failure Probability
\$100/day	0.1%
\$1/day	10%

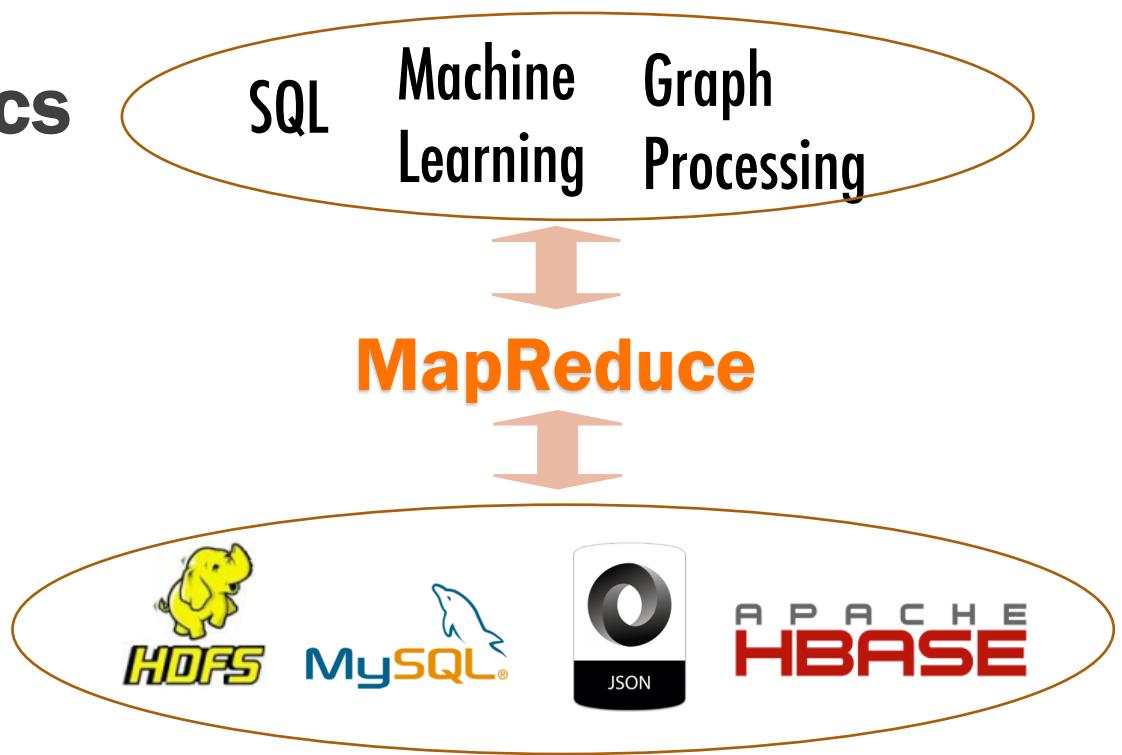
Reserved Instance      Spot Instance

The diagram illustrates the trade-off between the number of machines and failure probability, and how this relates to AWS instance types. On the left, a table shows that as the number of machines increases, the failure probability drops significantly. On the right, a table shows the cost and failure probability for different AWS instance types: Reserved Instances (cost \$100/day, 0.1% failure) and Spot Instances (cost \$1/day, 10% failure). Red arrows point from the 'Failure Probability' column of the machine table to the 'Failure Probability' row of the AWS table, indicating that more machines correspond to lower failure probabilities. A blue arrow points from the 'Cost for 5 nodes' column of the machine table to the 'Cost for 5 nodes' row of the AWS table, indicating that more machines correspond to higher costs.

# Why MapReduce?

---

## 2. Complex Analytics



## 3. Heterogeneous Storage Systems

---

# **The Great MapReduce Debate**

## **(2008-2010)**

<http://tiny.cc/mapreduce-debate>

# MapReduce vs. SQL

---

**Map** (k, v)



**SELECT** Map(v)  
**FROM** Table

**Reduce** (k, v\_list)



**SELECT** Reduce(v)  
**FROM** Table  
**GROUP BY** k

# MapReduce: A Major Step Backwards

---



1. MapReduce is a step backwards in database access
2. MapReduce is a poor implementation
3. MapReduce is not novel
4. MapReduce is missing features
5. MapReduce is incompatible with the DBMS tools

Dewitt, D. and Stonebraker, M. *MapReduce: A Major Step Backwards* blogpost; January 17, 2008

# Comments From The Other Side

---



VS.



**MapReduce** is a program model  
rather than a **database** system

From Stonebraker et al.

**MapReduce complements DBMSs since databases are not designed for extract-transform-load tasks, a MapReduce specialty.**

BY MICHAEL STONEBRAKER, DANIEL ABADI,  
DAVID J. DEWITT, SAM MADDEN, ERIK PAULSON,  
ANDREW PAVLO, AND ALEXANDER RASIN

# MapReduce and Parallel DBMSs: Friends or Foes?



From Dean and Ghemawat

**MapReduce advantages over parallel databases include storage-system independence and fine-grain fault tolerance for large jobs.**

BY JEFFREY DEAN AND SANJAY GHEMAWAT

# MapReduce: A Flexible Data Processing Tool

# What They Agree On?

---

Advantages of MapReduce:

1. Fault Tolerant
2. Complex Analytics
3. Heterogeneous Storage Systems
4. No Data Loading Requirement

Both should Learn from Each Other

# Who won the debate?

---



Nobody is writing  
MapReduce code right now.



Many new systems (e.g., Spark,  
HIVE) were built on MapReduce

# Summary

---

Why Relational Database? ( Mid 1970 – Now)

Why NoSQL? (Mid 2000 – Now)

Why MapReduce? (Mid 2000 – Now)