

# CMPT 354: Database System I

Lecture 4. SQL Advanced

# Announcements!

- A1 is due today
- A2 is released (due in 2 weeks)

# Outline

- **Joins**
  - Inner Join
  - Outer Join
- **Aggregation Queries**
  - Simple Aggregations
  - Group By
  - Having
- **Discussion**

# Joins: Recap

Student

| name | gpa |
|------|-----|
| Mary | 3.8 |
| Tom  | 3.6 |
| Jack | 3.7 |

Enroll

| stdName | course |
|---------|--------|
| Mary    | 354    |
| Tom     | 354    |
| Tom     | 454    |
| Bob     | 354    |

```
SELECT name, course  
FROM Student, Enroll  
WHERE name = stdName
```

name gpa

|      |     |
|------|-----|
| Mary | 354 |
| Tom  | 354 |
| Tom  | 454 |

# Two equivalent ways to write joins

```
SELECT name, course  
FROM Student, Enroll  
WHERE name = stdName
```



```
SELECT name, course  
FROM Student JOIN Enroll ON  
name = stdName
```

# Join Types

```
SELECT name, course  
FROM Student INNER JOIN Enroll ON  
      name = stdName
```

```
SELECT name, course  
FROM Student FULL OUTER JOIN Enroll ON  
      name = stdName
```

```
SELECT name  
FROM Student LEFT OUTER JOIN Enroll ON  
      name = stdName
```

```
SELECT name  
FROM Student RIGHT OUTER JOIN Enroll ON  
      name = stdName
```

# Join Types

```
SELECT name, course  
FROM Student           JOIN Enroll ON  
      name = stdName
```

```
SELECT name, course  
FROM Student FULL        JOIN Enroll ON  
      name = stdName
```

```
SELECT name  
FROM Student LEFT         JOIN Enroll ON  
      name = stdName
```

```
SELECT name  
FROM Student RIGHT        JOIN Enroll ON  
      name = stdName
```

# Left Join

Student

| name | gpa |
|------|-----|
| Mary | 3.8 |
| Tom  | 3.6 |
| Jack | 3.7 |

Enroll

| stdName | course |
|---------|--------|
| Mary    | 354    |
| Tom     | 354    |
| Tom     | 454    |
| Bob     | 354    |

```
SELECT name, course  
FROM Student LEFT JOIN Enroll ON  
name = stdName
```

We want to include all students no matter whether they enroll a course or not. How?

```
SELECT name, course  
FROM Student LEFT JOIN Enroll ON  
name = stdName
```

**Student**

| <b>name</b> | <b>gpa</b> |
|-------------|------------|
| Mary        | 3.8        |
| Tom         | 3.6        |
| Jack        | 3.7        |

**Enroll**

| <b>stdName</b> | <b>course</b> |
|----------------|---------------|
| Mary           | 354           |
| Tom            | 354           |
| Tom            | 454           |
| Bob            | 354           |

**Output**

| <b>name</b> | <b>course</b> |
|-------------|---------------|
| Mary        | 354           |
| Tom         | 354           |
| Tom         | 454           |
| Jack        | NULL          |

```
SELECT name, course  
FROM Student RIGHT JOIN Enroll ON  
name = stdName
```

**Student**

| <b>name</b> | <b>gpa</b> |
|-------------|------------|
| Mary        | 3.8        |
| Tom         | 3.6        |
| Jack        | 3.7        |

**Enroll**

| <b>stdName</b> | <b>course</b> |
|----------------|---------------|
| Mary           | 354           |
| Tom            | 354           |
| Tom            | 454           |
| Bob            | 354           |

**Output**

| <b>name</b> | <b>course</b> |
|-------------|---------------|
| Mary        | 354           |
| Tom         | 354           |
| Tom         | 454           |
| NULL        | 354           |

```
SELECT name, course  
FROM Enroll FULL JOIN Student ON  
name = stdName
```

Enroll

| stdName | course |
|---------|--------|
| Mary    | 354    |
| Tom     | 354    |
| Tom     | 454    |
| Bob     | 354    |

Student

| name | gpa |
|------|-----|
| Mary | 3.8 |
| Tom  | 3.6 |
| Jack | 3.7 |

Output

| name | course |
|------|--------|
| Mary | 354    |
| Tom  | 354    |
| Tom  | 454    |
| Jack | NULL   |
| NULL | 354    |

# Outer Join

TableA (**LEFT/RIGHT/FULL**) JOIN TableB

- Left outer join:
  - Include tuples from tableA even if no match
- Right outer join:
  - Include tuples from tableA even if no match
- Full outer join:
  - Include tuples from both even if no match

# Exercise - 1

```
SELECT name, course  
FROM Student LEFT JOIN Enroll ON  
name = stdName AND course = 354
```

Student

| name | gpa |
|------|-----|
| Mary | 3.8 |
| Tom  | 3.6 |
| Jack | 3.7 |

Enroll

| stdName | course |
|---------|--------|
| Mary    | 354    |
| Tom     | 354    |
| Tom     | 454    |
| Bob     | 354    |

| name | course |
|------|--------|
| Mary | 354    |
| Tom  | 354    |
| Jack | NULL   |

(A)

| name | course |
|------|--------|
| Mary | 354    |
| Tom  | 354    |

(B)

# Exercise - 2

```
SELECT name, course  
FROM Student LEFT JOIN Enroll ON  
name = stdName  
WHERE course = 354
```

Student

| name | gpa |
|------|-----|
| Mary | 3.8 |
| Tom  | 3.6 |
| Jack | 3.7 |

Enroll

| stdName | course |
|---------|--------|
| Mary    | 354    |
| Tom     | 354    |
| Tom     | 454    |
| Bob     | 354    |

| name | course |
|------|--------|
| Mary | 354    |
| Tom  | 354    |
| Jack | NULL   |

(A)

| name | course |
|------|--------|
| Mary | 354    |
| Tom  | 354    |

(B)

# Outline

- **Joins**
  - Inner Join
  - Outer Join
- **Aggregation Queries**
  - Simple Aggregations
  - Group By
  - Having
- **Discussion**

# Simple Aggregation

```
SELECT agg(column)
FROM   <table name>
WHERE  <conditions>
```

**agg** = COUNT, SUM, AVG, MAX, MIN, etc.

Except count, all aggregations apply to a single attribute

# Examples

| name  | gender | gpa |
|-------|--------|-----|
| Bob   | M      | 3   |
| Mike  | M      | 3   |
| Alice | F      | 3   |
| Mary  | F      | 4   |
| Tom   | M      | 4   |

```
SELECT COUNT(*) FROM Student
```

5

```
SELECT SUM(gpa) FROM Student
```

17

```
SELECT AVG(gpa) FROM Student
```

3.4

```
SELECT MIN(gpa) FROM Student
```

3

```
SELECT MAX(gpa) FROM Student
```

4

# Examples

| name  | gender | gpa |
|-------|--------|-----|
| Bob   | M      | 3   |
| Mike  | M      | 3   |
| Alice | F      | 3   |
| Mary  | F      | 4   |
| Tom   | M      | 4   |

```
SELECT COUNT(DISTINCT gpa) FROM Student
```

2

```
SELECT SUM(DISTINCT gpa) FROM Student
```

7

```
SELECT AVG(gpa) FROM Student  
WHERE gender = 'Female'
```

3.5

# The need for Group By

- How to get AVG(gpa) for each gender?

```
SELECT AVG(gpa) FROM Student WHERE gender = 'M'
```

```
SELECT AVG(gpa) FROM Student WHERE gender = 'F'
```

- How to get AVG(gpa) for each age?

```
SELECT AVG(gpa) FROM Student WHERE age = 18
```

```
SELECT AVG(gpa) FROM Student WHERE age = 19
```

```
SELECT AVG(gpa) FROM Student WHERE age = 20
```

•  
•  
•

# Grouping and Aggregation

```
SELECT agg(column)
FROM   <table name>
WHERE  <conditions>
GROUP BY <columns>
```

- How to get AVG(gpa) for each gender?

```
SELECT AVG(gpa) FROM Student GROUP BY gender
```

- How to get AVG(gpa) for each age?

```
SELECT AVG(gpa) FROM Student GROUP BY age
```

# Grouping and Aggregation

- How is the following query processed?

```
SELECT gender, AVG(gpa)
FROM Student
WHERE gpa > 2.5
GROUP BY gender
```

- Semantics of the query
  1. Compute the **FROM** and **WHERE** clauses
  2. Group by the attributes in the **GROUP BY**
  3. Compute the **SELECT** clause: grouped attributes and aggregates

# 1. Compute the **FROM** and **WHERE** clauses

```
SELECT gender, AVG(gpa)
  FROM Student
 WHERE gpa > 2.5
 GROUP BY gender
```

| name  | gender | gpa |
|-------|--------|-----|
| Bob   | M      | 2   |
| Mike  | M      | 3   |
| Alice | F      | 3   |
| Mary  | F      | 4   |
| Tom   | M      | 3   |



| name  | gender | gpa |
|-------|--------|-----|
| Mike  | M      | 3   |
| Alice | F      | 3   |
| Mary  | F      | 4   |
| Tom   | M      | 3   |

## 2. Group by the attributes in the GROUP BY

```
SELECT gender, AVG(gpa)
FROM Student
WHERE gpa > 2.5
GROUP BY gender
```

| name  | gender | gpa |
|-------|--------|-----|
| Mike  | M      | 3   |
| Alice | F      | 3   |
| Mary  | F      | 4   |
| Tom   | M      | 3   |



| gender | name  | gpa |
|--------|-------|-----|
| M      | Mike  | 3   |
|        | Tom   | 3   |
| F      | Alice | 3   |
|        | Mary  | 4   |

### 3. Compute the **SELECT** clause: grouped attributes and aggregates

```
SELECT gender, AVG(gpa)
```

```
FROM Student
```

```
WHERE gpa > 2.5
```

```
GROUP BY gender
```

| gender | name  | gpa |
|--------|-------|-----|
| M      | Mike  | 3   |
|        | Tom   | 3   |
| F      | Alice | 3   |
|        | Mary  | 4   |



| gender | AVG(gpa) |
|--------|----------|
| M      | 3        |
| F      | 3.5      |

# Exercise: Empty Group

| name  | gender | gpa |
|-------|--------|-----|
| Bob   | M      | 3   |
| Mike  | M      | 3   |
| Alice | F      | 4   |
| Mary  | F      | 4   |
| Tom   | M      | 3   |

```
SELECT gender, AVG(gpa)  
FROM Student  
WHERE gpa > 3.5  
GROUP BY gender
```

| gender | AVG(gpa) |
|--------|----------|
| F      | 4        |

VS

| gender | AVG(gpa) |
|--------|----------|
| F      | 4        |
| M      | NULL     |

(A)

(B)

# Exercise: Empty Group

| name  | gender | gpa |
|-------|--------|-----|
| Bob   | M      | 3   |
| Mike  | M      | 3   |
| Alice | F      | 4   |
| Mary  | F      | 4   |
| Tom   | M      | 3   |



```
SELECT gender, AVG(gpa)
FROM Student
WHERE gpa > 3.5
GROUP BY gender
```

| name  | gender | gpa |
|-------|--------|-----|
| Alice | F      | 4   |
| Mary  | F      | 4   |

# Exercise: Empty Group

| name  | gender | gpa |
|-------|--------|-----|
| Bob   | M      | 3   |
| Mike  | M      | 3   |
| Alice | F      | 4   |
| Mary  | F      | 4   |
| Tom   | M      | 3   |



```
SELECT gender, AVG(gpa)
FROM Student
WHERE gpa > 3.5
GROUP BY gender
```

| name  | gender | gpa |
|-------|--------|-----|
| Alice | F      | 4   |
| Mary  | F      | 4   |



| gender | name  | gpa |
|--------|-------|-----|
| F      | Alice | 4   |
|        | Mary  | 4   |

# Exercise: Empty Group

| name  | gender | gpa |
|-------|--------|-----|
| Bob   | M      | 3   |
| Mike  | M      | 3   |
| Alice | F      | 4   |
| Mary  | F      | 4   |
| Tom   | M      | 3   |

```
SELECT gender, AVG(gpa)  
FROM Student  
WHERE gpa > 3.5  
GROUP BY gender
```



| name  | gender | gpa |
|-------|--------|-----|
| Alice | F      | 4   |
| Mary  | F      | 4   |



| gender | name  | gpa |
|--------|-------|-----|
| F      | Alice | 4   |
|        | Mary  | 4   |



| gender | AVG(gpa) |
|--------|----------|
| F      | 4        |

# Exercise: Invalid Selection

| name  | gender | gpa |
|-------|--------|-----|
| Bob   | M      | 3   |
| Mike  | M      | 3   |
| Alice | F      | 4   |
| Mary  | F      | 4   |
| Tom   | M      | 3   |

```
SELECT gender, AVG(gpa), name  
FROM Student  
WHERE gpa > 3.5  
GROUP BY gender
```

| gender | AVG(gpa) | name  |
|--------|----------|-------|
| F      | 4        | Alice |

VS

| gender | AVG(gpa) | name |
|--------|----------|------|
| F      | 4        | Mary |

(A)

(B)

# Exercise: Invalid Selection

| name  | gender | gpa |
|-------|--------|-----|
| Bob   | M      | 3   |
| Mike  | M      | 3   |
| Alice | F      | 4   |
| Mary  | F      | 4   |
| Tom   | M      | 3   |



```
SELECT gender, AVG(gpa), name  
FROM Student  
WHERE gpa > 3.5  
GROUP BY gender
```

| name  | gender | gpa |
|-------|--------|-----|
| Alice | F      | 4   |
| Mary  | F      | 4   |

# Exercise: Invalid Selection

| name  | gender | gpa |
|-------|--------|-----|
| Bob   | M      | 3   |
| Mike  | M      | 3   |
| Alice | F      | 4   |
| Mary  | F      | 4   |
| Tom   | M      | 3   |



```
SELECT gender, AVG(gpa), name  
FROM Student  
WHERE gpa > 3.5  
GROUP BY gender
```

| name  | gender | gpa |
|-------|--------|-----|
| Alice | F      | 4   |
| Mary  | F      | 4   |



| gender | name  | gpa |
|--------|-------|-----|
| F      | Alice | 4   |
|        | Mary  | 4   |

# Exercise: Invalid Selection

Everything in SELECT must be either a GROUP-BY attribute, or an aggregate

| name  | gender | gpa |
|-------|--------|-----|
| Bob   | M      | 3   |
| Mike  | M      | 3   |
| Alice | F      | 4   |
| Mary  | F      | 4   |
| Tom   | M      | 3   |

```
SELECT gender, AVG(gpa), name  
FROM Student  
WHERE gpa > 3.5  
GROUP BY gender
```



| name  | gender | gpa |
|-------|--------|-----|
| Alice | F      | 4   |
| Mary  | F      | 4   |



| gender | name  | gpa |
|--------|-------|-----|
| F      | Alice | 4   |
|        | Mary  | 4   |



| gender | AVG(gpa) | name |
|--------|----------|------|
| F      | 4        | ???  |
|        |          |      |

# HAVING Clause

- Specify which groups you are interested in

```
SELECT agg(column)
FROM   <table name>
WHERE  <conditions>
GROUP BY <columns>
HAVING <columns>
```

# HAVING Clause

- Same query as before, except that we require each group has more than 10 students

```
SELECT AVG(gpa), gender  
FROM Student  
WHERE gpa > 2.5  
GROUP BY gender  
HAVING COUNT(*) > 10
```

HAVING clause contains conditions on aggregates.

# Order of Evaluation

|          |                   |
|----------|-------------------|
| SELECT   | S                 |
| FROM     | $R_1, \dots, R_n$ |
| WHERE    | $C_1$             |
| GROUP BY | $a_1, \dots, a_k$ |
| HAVING   | $C_2$             |

- Create the cross product of the tables in the **FROM** clause
- Remove rows not meeting the **WHERE** condition
- Divide records into groups by the **GROUP BY** clause
- Remove groups not meeting the **HAVING** clause
- Create one row for each group and remove columns not in the **SELECT** clause

# Exercise

StudentInfo

| name  | gender | gpa |
|-------|--------|-----|
| Bob   | M      | 3   |
| Mike  | M      | 3   |
| Alice | F      | 4   |
| Mary  | F      | 4   |
| Tom   | M      | 3   |

```
SELECT gender, AVG(gpa)
FROM StudentInfo
WHERE gpa > 2.5
GROUP BY gender
HAVING COUNT(*) > 2
```

```
SELECT gender, AVG(gpa)
FROM StudentInfo
WHERE gpa > 2.5
GROUP BY gender
HAVING SUM(gpa) < 9
```

| gender | AVG(gpa) |
|--------|----------|
| M      | 3        |

(A)

| gender | AVG(gpa) |
|--------|----------|
| F      | 4        |

(B)

| gender | AVG(gpa) |
|--------|----------|
| M      | 3        |
| F      | 4        |

(C)

# Discussion

Customer = {customerID, firstName, lastName, income, birthDate}

Account = {accNumber, type, balance, branchNumber<sup>FK-Branch</sup>}

Owns = {customerID<sup>FK-Customer</sup>, accNumber<sup>FK-Account</sup>}

Q1. Show me the name of the richest customer?

Q2. Show me the customers who have at least 2 accounts

# Discussion

---

Employee = {sin, firstName, lastName, salary, branchNumber<sup>FK-Branch</sup>}

Branch = {branchNumber, branchName, managerSIN<sup>FK-Employee</sup>, budget}

Q3. Show me the names of branch managers?

Q4. Show me the average salary of the employees in each branch

# Acknowledge

- Some lecture slides were copied from or inspired by the following course materials
  - “W4111: Introduction to databases” by Eugene Wu at Columbia University
  - “CSE344: Introduction to Data Management” by Dan Suciu at University of Washington
  - “CMPT354: Database System I” by John Edgar at Simon Fraser University
  - “CS186: Introduction to Database Systems” by Joe Hellerstein at UC Berkeley
  - “CS145: Introduction to Databases” by Peter Bailis at Stanford
  - “CS 348: Introduction to Database Management” by Grant Weddell at University of Waterloo