

CMPT 354:

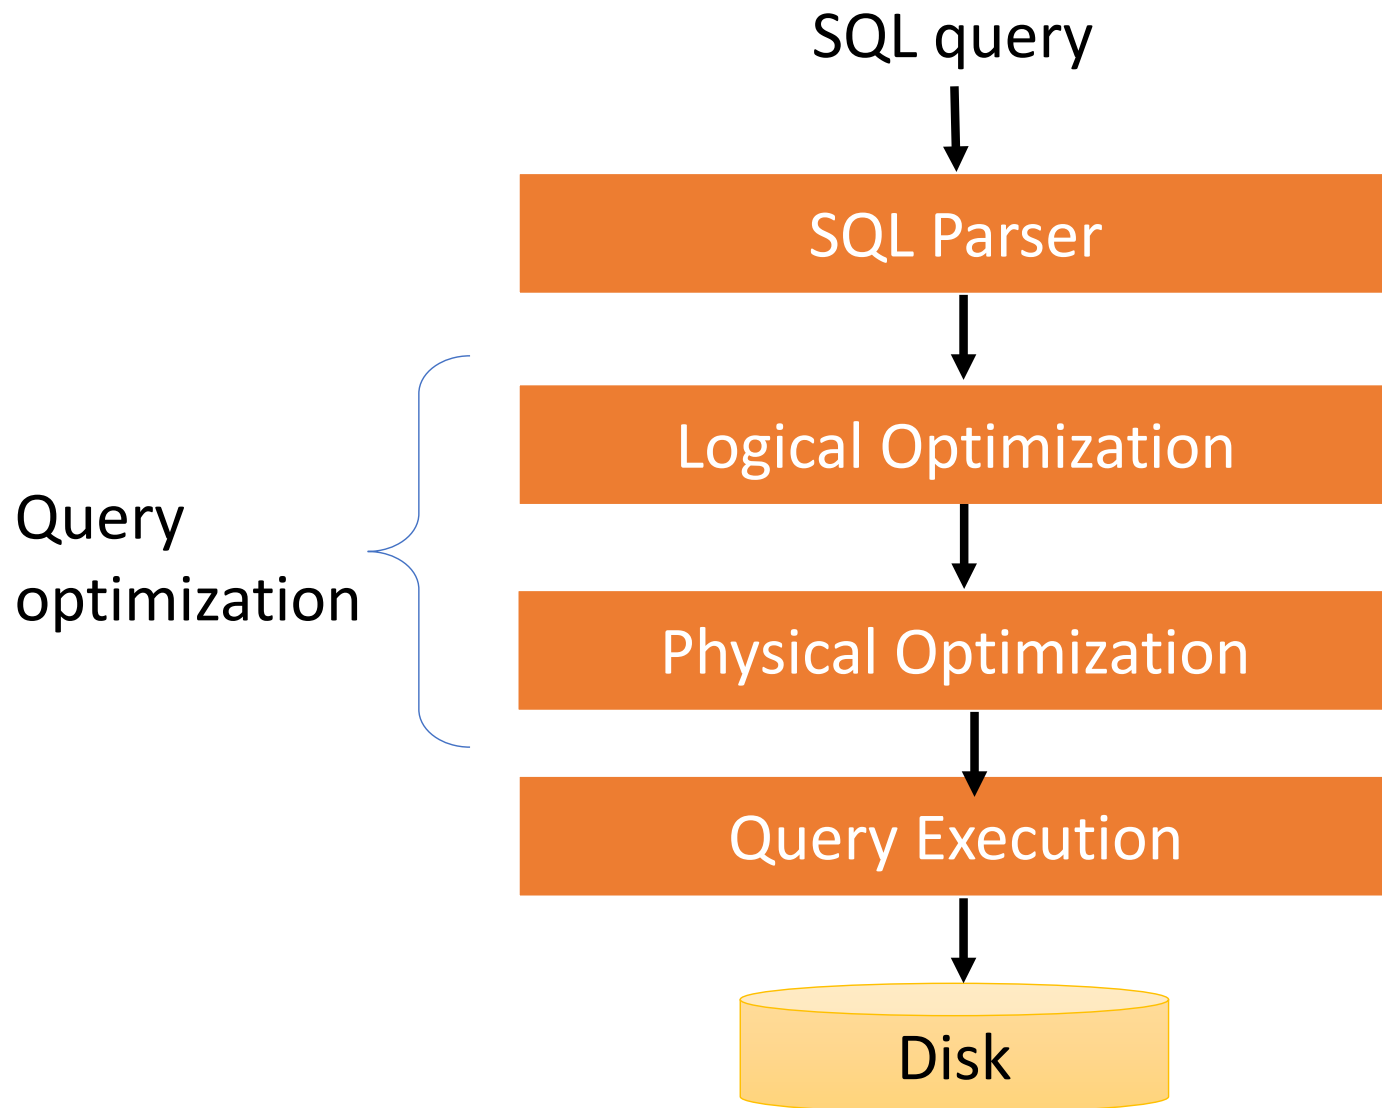
Database System I

Lecture 6. Basics of Query Processing and Indexing

Outline

- Query Processing
 - What happens when an SQL query is issued?
- Indexing
 - How to speed up query performance?

Query Processing Steps



Example

- **Offering** (oID, dept, cNum, term, instructor)
- **Took** (sID, oID, grade)

Q: Student number of all students who have taken CMPT 354

```
SELECT sID
FROM   Offering O, Took T
WHERE  O.oID = T.oID
       AND O.dept = 'CMPT'
       AND O.cNum = '354'
```

Offering (oID, dept, cNum, term, instructor)

Took (sID, oID, grade)

SQL Parser


- From the input SQL text to a logical plan

```
SELECT sID
FROM   Offering O, Took T
WHERE  O.oID = T.oID
      AND O.dept = 'CMPT'
      AND O.cNum = '354'
```



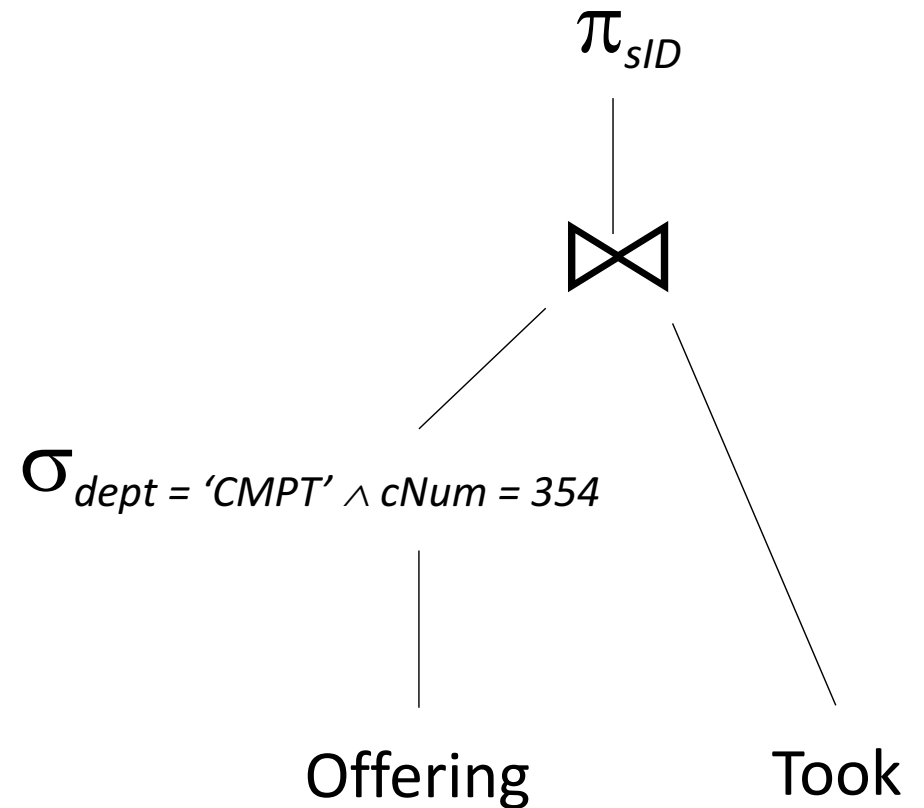

$\pi_{sID} (\sigma_{dept = 'CMPT' \wedge cNum = 354} (Offering \bowtie Took))$

Relational algebra expression is also called the “**logical query plan**”



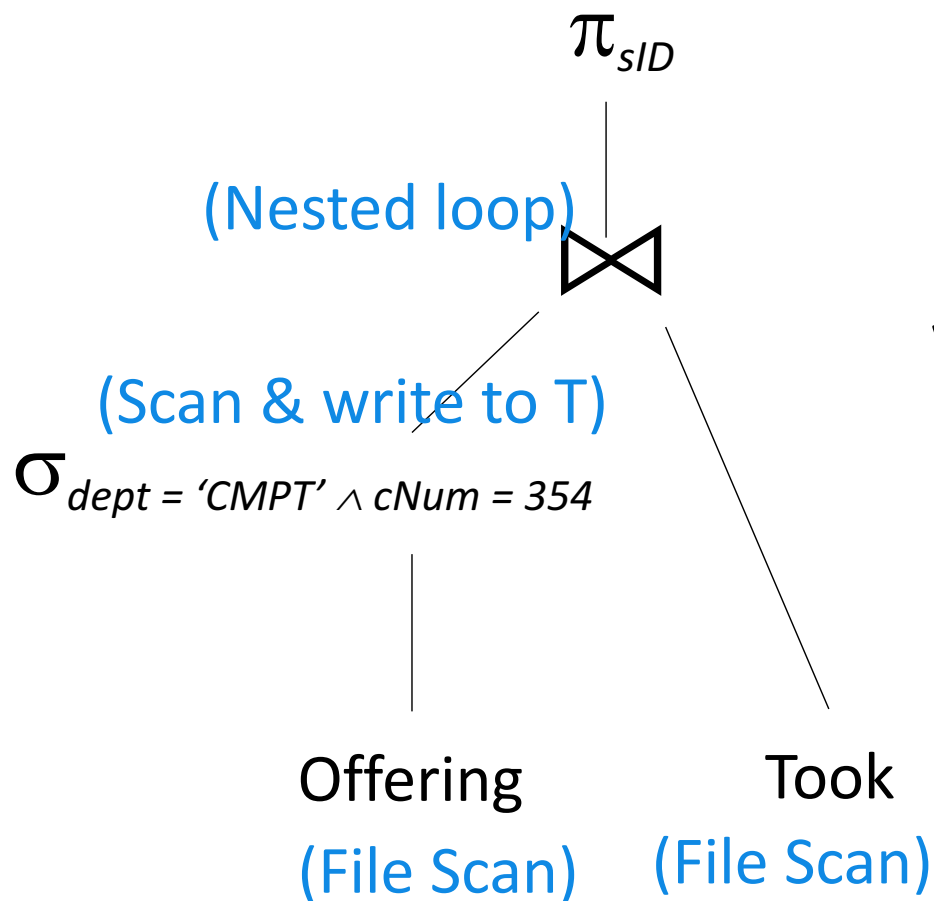
Logical Optimization

- Find the **optimal** logical plan

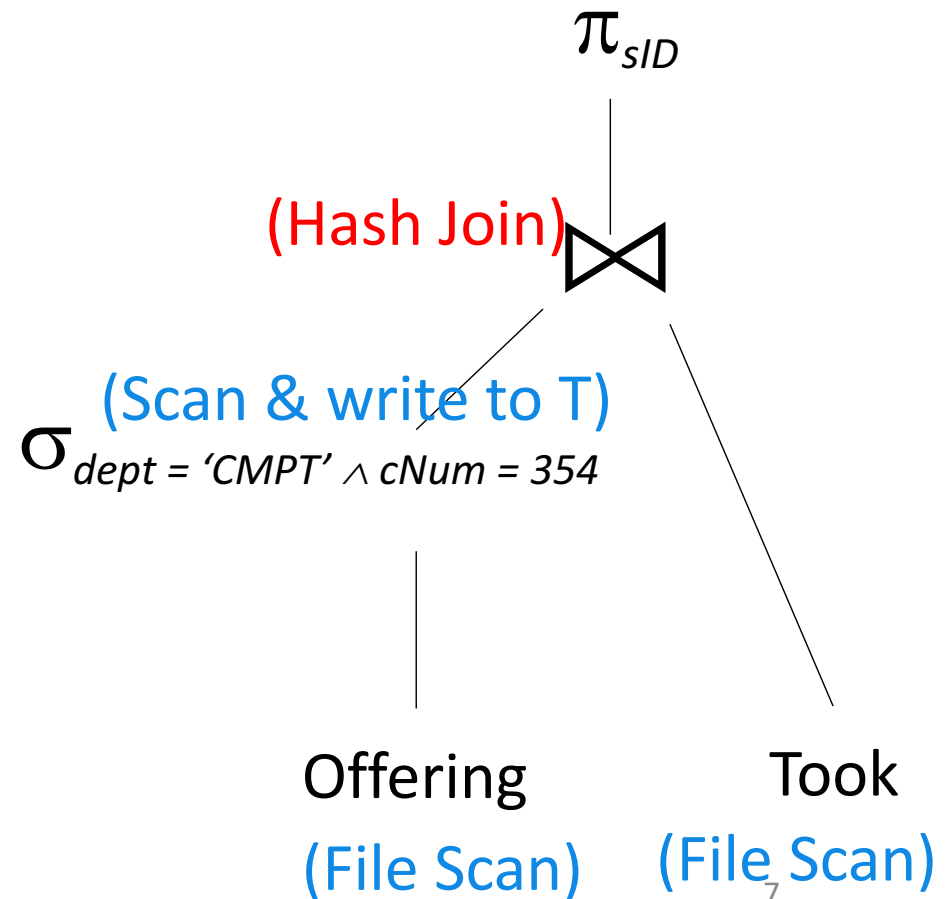


Physical Optimization

- Find the **optimal physical plan**

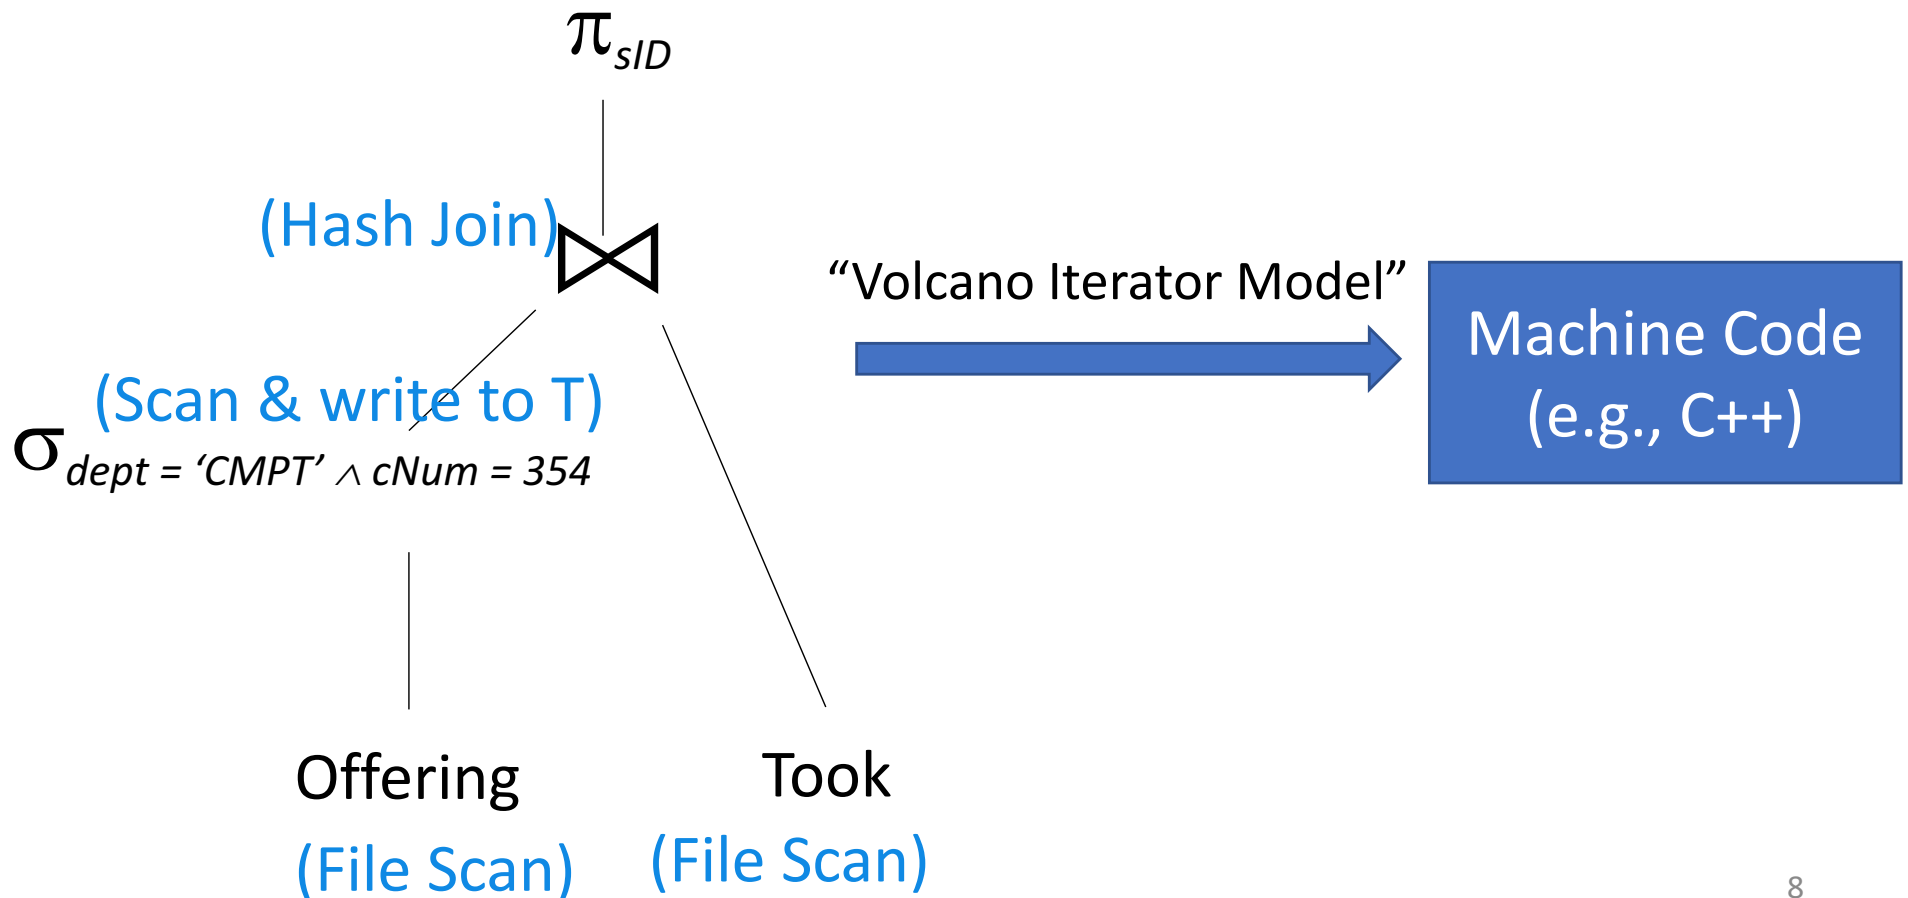


V.S.



Query Execution

- From a physical plan to actual machine code



Summary

- **Logical plans:**
 - Created by the parser from the input SQL text
 - Expressed as a relational algebra tree
 - Each SQL query has many possible logical plans
- **Physical plans:**
 - Goal is to choose an efficient implementation for each operator in the RA
 - Each logical plan has many possible physical plans
- **Query Optimization:**
 - Find the optimal logical plan
 - Find the optimal physical plan

Outline

- Query Processing
 - What happens when an SQL query is issued?
- Indexing
 - How to speed up query performance?

Query Performance

- My database application is too slow... why?
- One of the queries is very slow... why?
- To address these problems, we need to understand:
 - How is data organized on disk
 - What is an index
 - How to select indexes

Data Storage

- DBMSs store data in files
- Most common organization is row-wise storage
- On disk, a file is split into **blocks**
- Each block contains a set of tuples

sID	dept	cNum	Term	instructor
10	CMPT	345	SP 2018	Jiannan
20	CMPT	454	FA 2018	Martin
...

10	CMPT	345	SP 2018	Jiannan
20	CMPT	454	FA 2018	Martin

Block 1

30
40	...			

Block 2

50				
60				

Block 3

70				
80				

Block 4

In the example, we have 4 blocks with 2 tuples each

Scanning a Data File

- Data file is stored on Disk
- Consequence: Sequential IO is MUCH FASTER than random IO
 - Good: read blocks 1, 2, 3, 4, 5
 - Bad: read blocks 2342, 11, 321, 9
- Rule of thumb:
 - Random reading 1-2% of the file \approx sequential scanning the entire file

Data File Types

- **Heap file**
 - Unsorted
- **Sequential file**
 - Sorted according to some attribute(s) called key

Note: key here means something different from primary key: it just means that we order the file according to that attribute. In our example we ordered by **sID**. Might as well order by **instructor**, if that seems a better idea for the applications running on our database.

Index Motivation


Student(name, age)

- Suppose we want to search for students of a specific age
- ***First idea:*** Sort the records by age... we know how to do this fast!
- How many IO operations to search over ***N sorted*** records?
 - Simple scan: $O(N)$
 - Binary search: $O(\log_2 N)$

Could we get even cheaper search? E.g. go from $\log_2 N$
 $\rightarrow \log_{200} N$

Index Motivation

- What about if we want to **insert** a new student, but keep the list sorted?



- We would have to potentially shift **N** records, requiring up to $\sim 2 \cdot N/P$ IO operations (where $P = \#$ of records per page)!

Could we get faster insertions?

Index Motivation

- What about if we want to be able to search quickly along multiple attributes (e.g. not just age)?
 - We could keep multiple copies of the records, each sorted by one attribute set... this would take a lot of space

Can we get fast search over multiple attribute sets without taking too much space?

We'll create separate data structures called *indexes* to address all these points

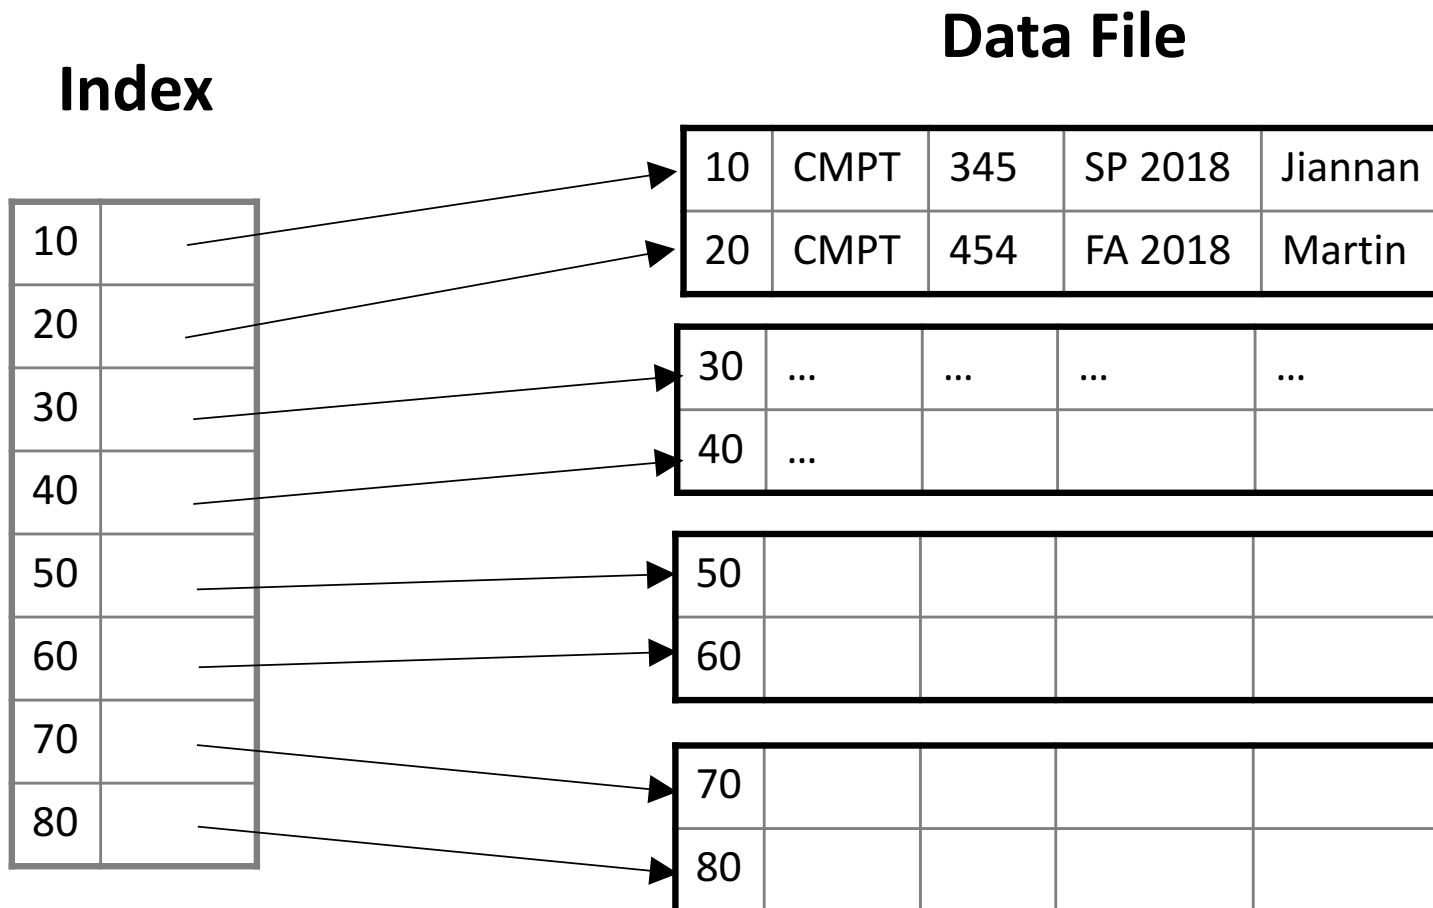
Index

- An additional file, that allows fast access to records in the data file given a search key
- The index contains (key, value) pairs:
 - The key = an attribute value (e.g., student ID or name)
 - The value = a pointer to the record
- An index can store the full rows it points to (*primary index*) or pointers to those rows (*secondary index*)
 - We'll mainly consider secondary indexes
- Could have many indexes for one table

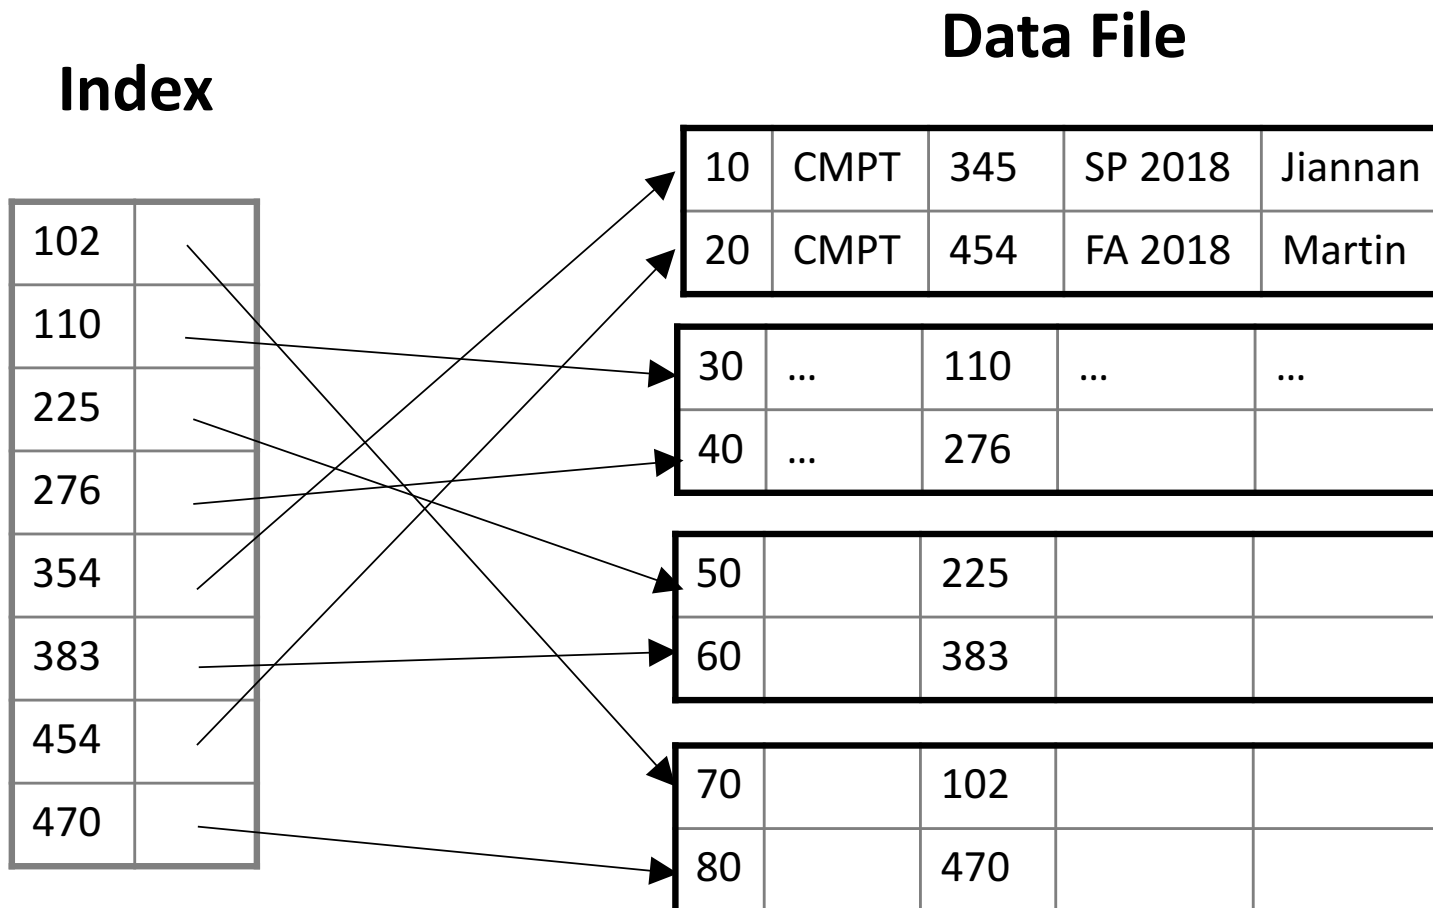
Different Keys

- **Primary key**
 - uniquely identifies a tuple
- **Key of the sequential file**
 - how the data file is sorted
- **Index key**
 - how the index is organized

Example 1: Index on sID



Example 2: Index on cNum



Index Organization

- Common indexes:
 - Hash tables
 - B+ trees
- Specialized indexes
 - R-trees
 - Inverted index
 - ...

B+ Tree Example

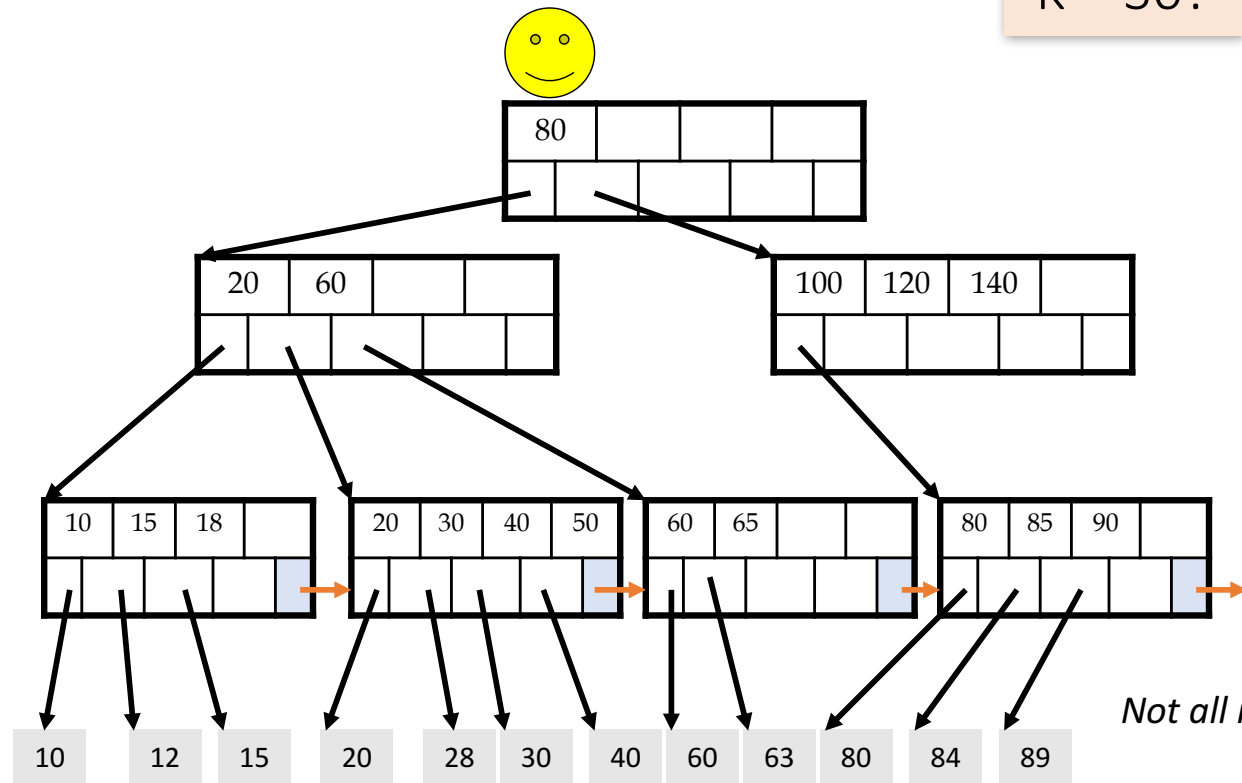
K = 30?

30 < 80

30 in [20,60)


30 in [30,40)

To the data!



Not all nodes pictured


Clustered vs. Unclustered Index



Clustered

Index File

Data file

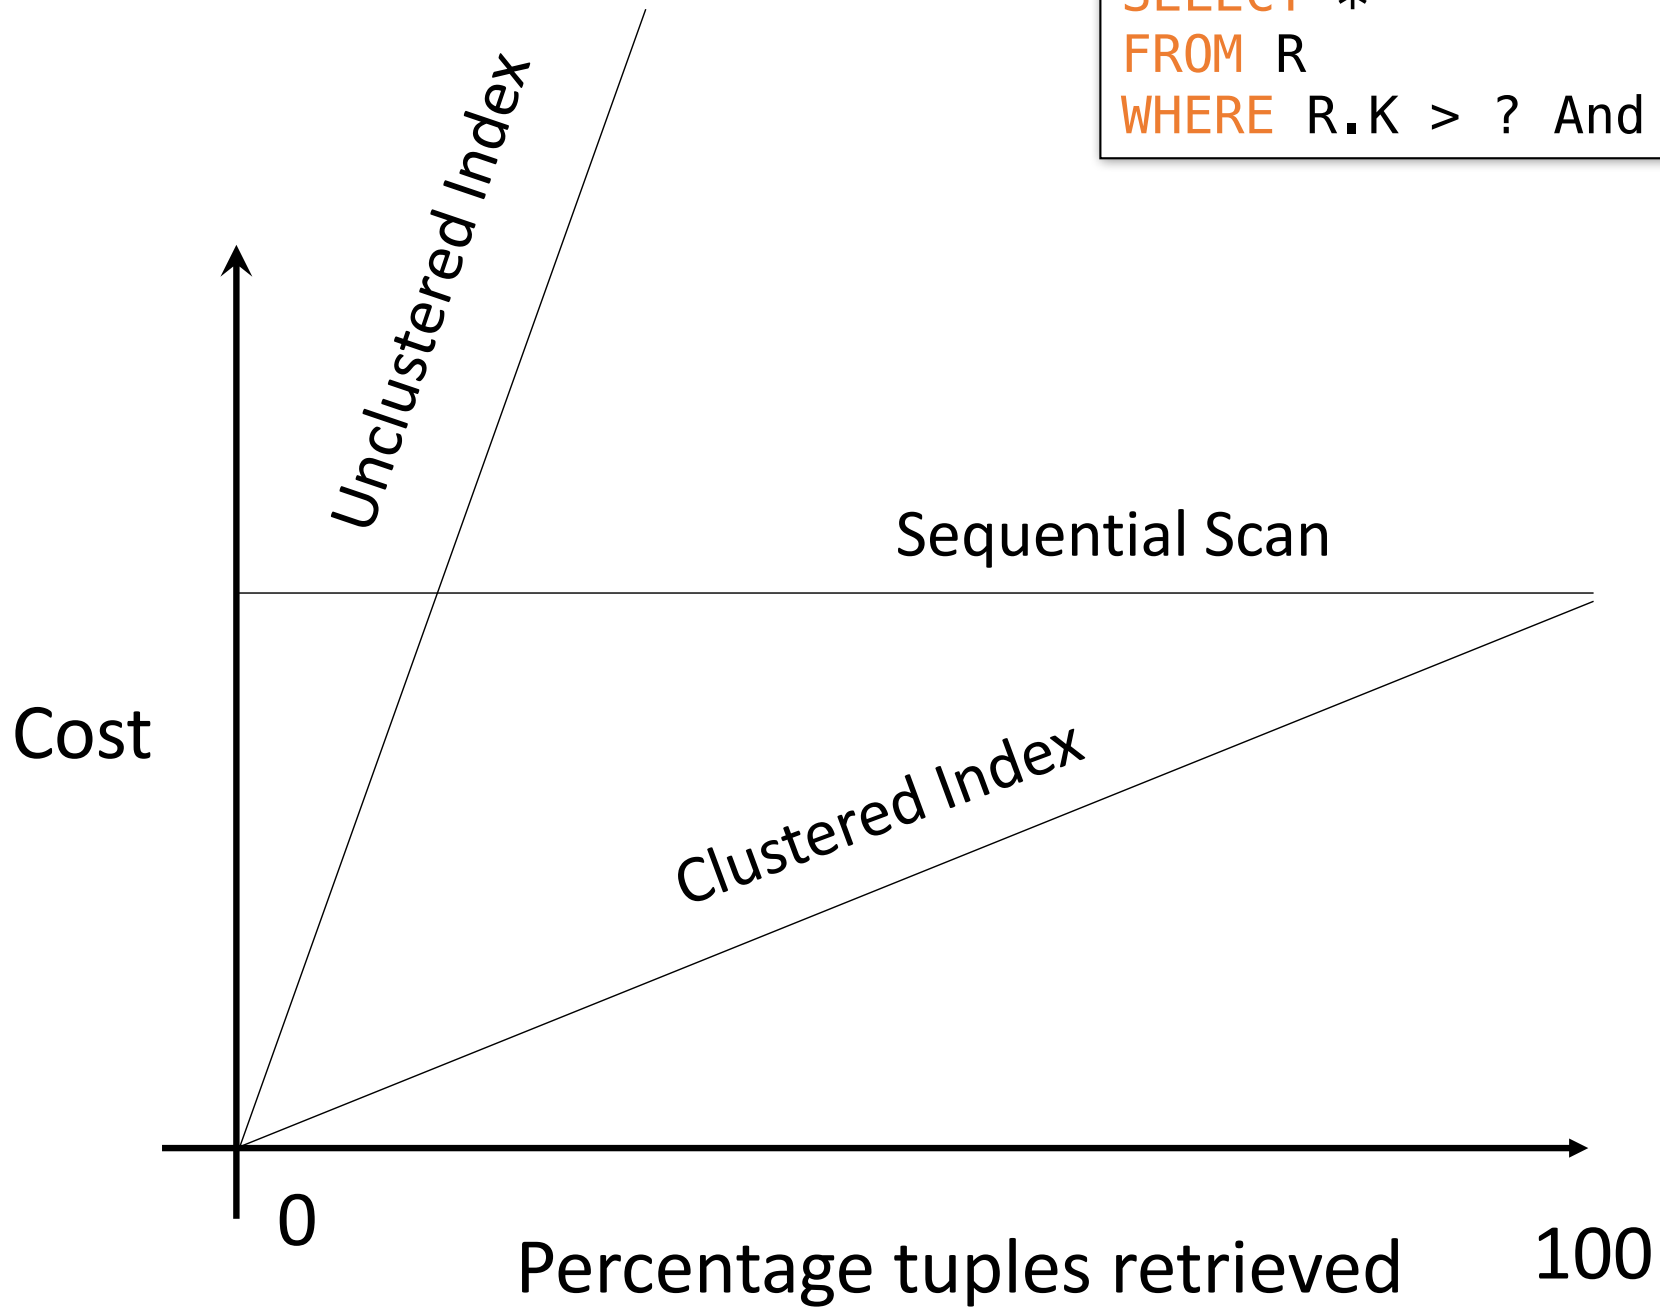


Unclustered

Clustered vs. Unclustered Index

- Recall that for a disk with block access, **sequential IO is much faster than random IO**
- For exact search, no difference between clustered / unclustered
- For range search over R values: difference between **1 random IO + R sequential IO**, and **R random IO**

```
SELECT *  
FROM R  
WHERE R.K > ? And R.K < ?
```



Summary

- **Logical plans:**
 - Created by the parser from the input SQL text
 - Expressed as a relational algebra tree
 - Each SQL query has many possible logical plans
- **Physical plans:**
 - Goal is to choose an efficient implementation for each operator in the RA
 - Each logical plan has many possible physical plans
- **Query Optimization:**
 - Find the optimal logical plan
 - Find the optimal physical plan

Summary

- Index = a file that enables direct access to records in another data file
 - B+ tree / Hash table
 - Clustered/unclustered
- Data resides on disk
 - Organized in blocks
 - Sequential IO is more efficient than random IO
 - Random read 1-2% of data worse than sequential scan of the entire file

Creating Indexes in SQL

- **Offering** (oID, dept, cNum, term, instructor)

```
CREATE INDEX IDX1 ON Offering(dept)
```

Which query(s) could be affected by IDX1?

(A)

```
SELECT oID FROM Offering  
WHERE dept = 'CMPT'
```

(B)

```
SELECT oID FROM Offering  
WHERE cNum = '354'
```

(C)

```
SELECT oID FROM Offering  
WHERE dept = 'CMPT' AND cNum = '354'
```

Creating Indexes in SQL

- Offering (oID, dept, cNum, term, instructor)

```
CREATE INDEX IDX2 ON Offering(dept, cNum)
```

Which query(s) could be affected by IDX2?

(A)

```
SELECT oID FROM Offering  
WHERE dept = 'CMPT'
```

(B)

```
SELECT oID FROM Offering  
WHERE cNum = '354'
```

(C)

```
SELECT oID FROM Offering  
WHERE dept = 'CMPT' AND cNum = '354'
```

Which Indexes?

- How many indexes could we create?
- Which indexes should we create?

Which Indexes?

- The index selection problem
 - Given a table, and a “workload” (SFU CourSys application with lots of SQL queries), decide which indexes to create (and which ones NOT to create!)
- Who does index selection:
 - The database administrator DBA
 - Semi-automatically, using a database administration tool

Index Selection: Which Search Key

- Make some attribute K a search key if the WHERE clause contains:
 - An exact match on K
 - A range predicate on K
 - A join on K

The Index Selection Problem 1

- Your workload is

100000 queries

```
SELECT sID  
FROM Student  
WHERE name = ?
```

100000 queries

```
SELECT sID  
FROM Student  
WHERE gender = ?
```

Which one is better?

- A. Index on name
- B. Index on gender

The Index Selection Problem 2

- Your workload is

100000 queries

```
SELECT sID  
FROM Student  
WHERE name like ?
```

100000 queries

```
SELECT sID  
FROM Student  
WHERE age = ?
```

Which one is better?

- A. Index on name
- B. Index on age

The Index Selection Problem 3

- Your workload is

100000 queries

```
SELECT sID  
FROM Student  
WHERE name = ?
```

100 queries

```
SELECT sID  
FROM Student  
WHERE age = ?
```

Which one(s) are useful?

- A. Index on name
- B. Index on age
- C. Index on name, age
- D. Index on age, name

The Index Selection Problem 4

- Your workload is

100000 queries

```
SELECT sID  
FROM Student  
WHERE fname = ?
```

100000 queries

```
SELECT sID  
FROM Student  
WHERE fname = ? AND age > ?
```

Which one is better?

- A. Index on (fname, age)
- B. Index on (age, fname)

The Index Selection Problem 5

- Your workload:

100000 queries

```
SELECT sID
FROM Student
WHERE name = ?
```

100 queries

```
SELECT sID
FROM Student
WHERE age = ?
```

100000 queries

```
INSERT INTO Student
VALUES (?, ..., ?)
```

Which one(s) are useful?

- A. Index on name
- B. Index on age
- C. Index on name, age
- D. Index on age, name

Basic Index Selection Guidelines

- Consider queries in workload in order of importance
- Consider relations accessed by query
 - No point indexing other relations
- Look at WHERE clause for possible search key
- Try to choose indexes that speed up multiple queries

Summary

- Query Processing
 - SQL Parser
 - Logical Optimization
 - Physical Optimization
 - Query Execution
- Indexing
 - Data Storage
 - Index motivation
 - Index Selection

Acknowledge

- Some lecture slides were copied from or inspired by the following course materials
 - “W4111: Introduction to databases” by Eugene Wu at Columbia University
 - “CSE344: Introduction to Data Management” by Dan Suciu at University of Washington
 - “CMPT354: Database System I” by John Edgar at Simon Fraser University
 - “CS186: Introduction to Database Systems” by Joe Hellerstein at UC Berkeley
 - “CS145: Introduction to Databases” by Peter Bailis at Stanford
 - “CS 348: Introduction to Database Management” by Grant Weddell at University of Waterloo