

CMPT 354: Database System I

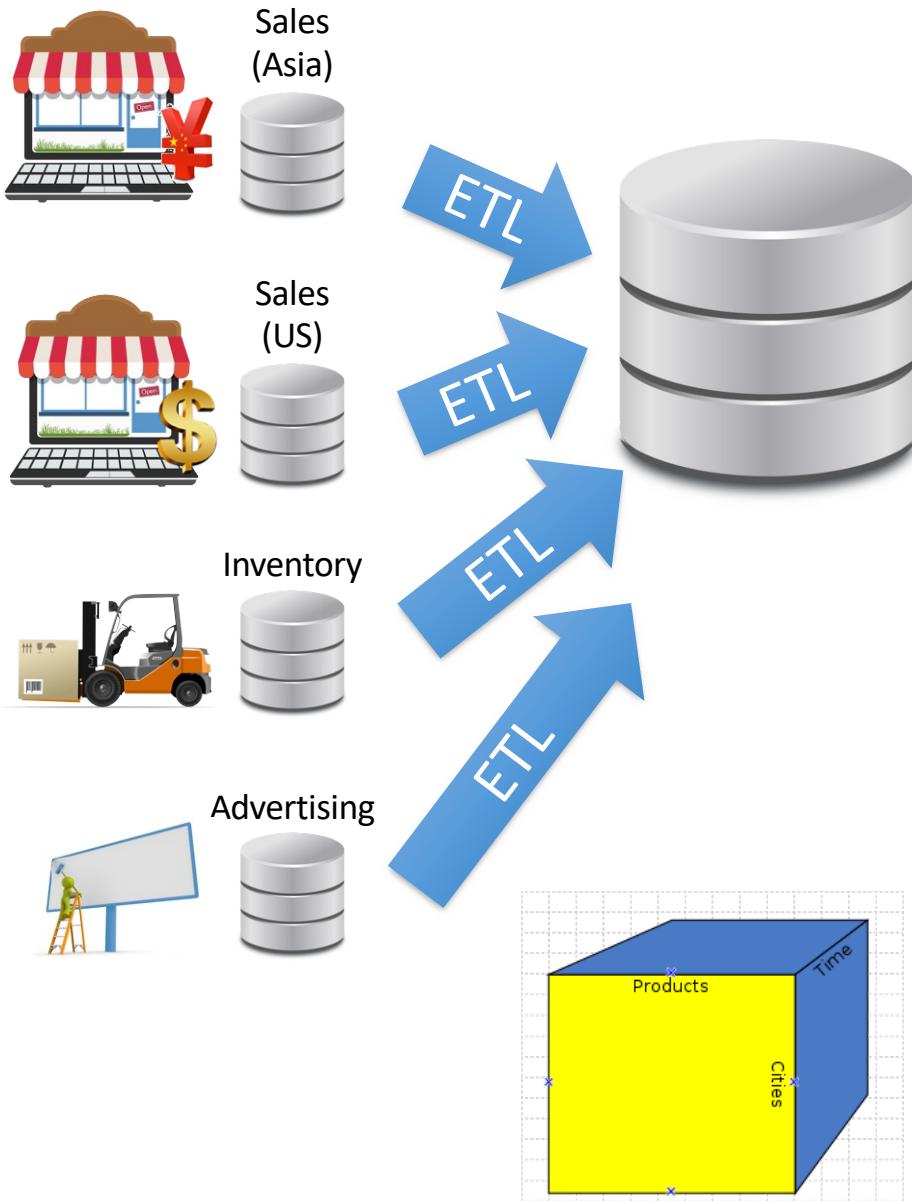
Lecture 12. Where to go from here

Where are we now

- **How to query a database**
 - SQL (Lec 3-4)
- **How to design a database**
 - ER Model (Lec 8)
 - Design Theory (Lec 9)
 - Relational Model (Lec 2)
- **How to develop a database application**
 - Database Application (Lec 10)
 - Transactions (Lec 12)
- **How to process a SQL query**
 - Relational Algebra (Lec 5)
 - Query Processing (Lec 6)
 - Query Optimization (Lec 12)

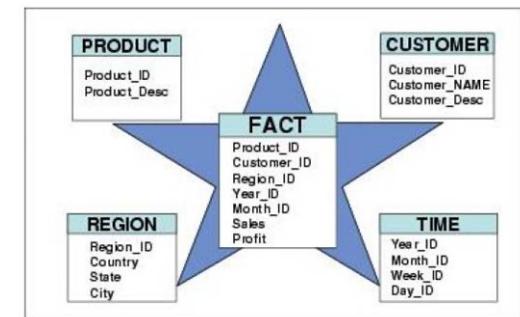
Where to go from here

- Data Warehouse
- Parallel Databases
- Big Data Processing
- Cloud Databases
- Data Science



Data Warehouse

Collects and organizes historical data from multiple sources



So far ...

- Star Schemas
- Data cubes



Data Warehouse

Collects and organizes historical data from multiple sources

- How do we deal with semi-structured and unstructured data?
- Do we really want to force a schema on load?

Data Warehouse

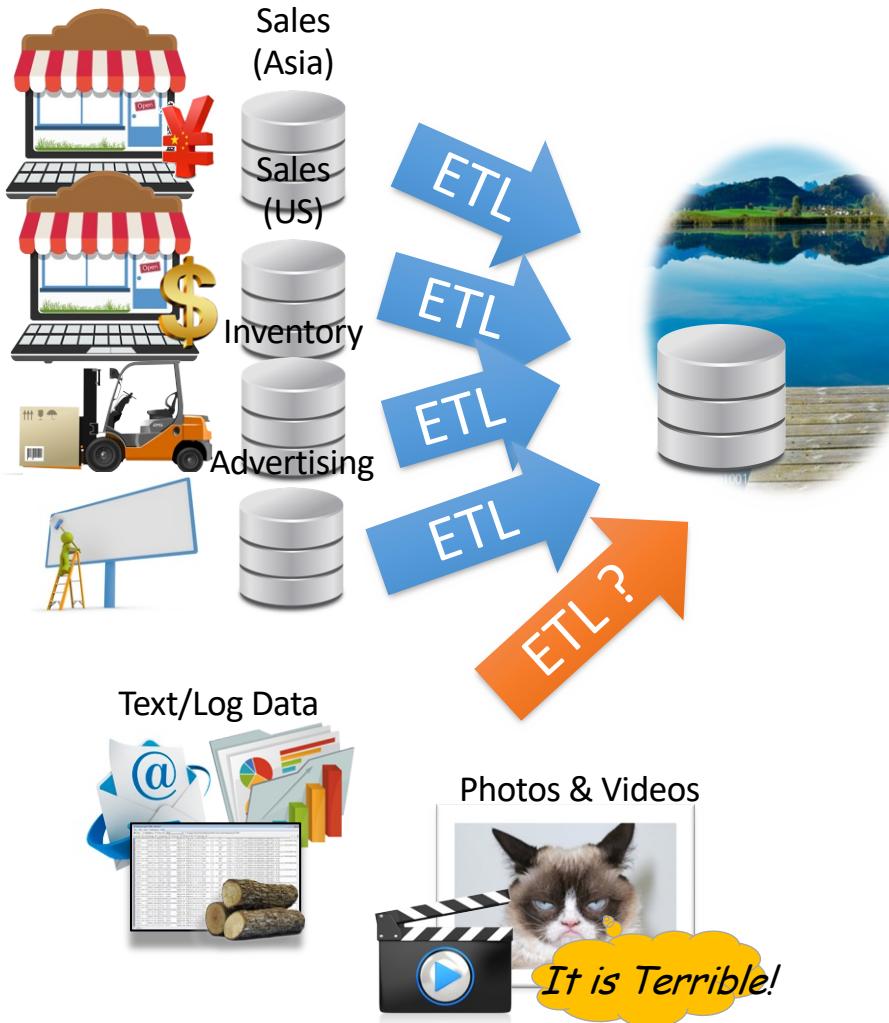


Collects and organizes historical data from multiple sources

How do we **load** and **process** this data in a relational system?

Do we read on load?

Depends on use ...
Can be difficult ...
Requires thought ...



Data Lake *

Store a copy of all the data

- in one place
- in its original “natural” form

Enable data consumers to choose how to transform and use data.

- *Schema on Read*

Enabled by new Tools:
Map-Reduce & Distributed Filesystems

What could go wrong?

* Still being defined...[Buzzword Disclaimer]

The Dark Side of Data Lakes



- Cultural shift: *Curate* → *Save Everything!*
 - Noise begins to dominate signal
- Limited data governance and planning
 - **Example:**
hdfs://important/joseph_big_file3.csv_with_json
 - **What** does it contain?
 - **When** and **who** created it?
 - No cleaning and verification → lots of dirty data
 - New tools are more complex and old tools no longer work

Enter the data engineer

A Brighter Future for Data Lakes

Enter the data engineer

- Data engineers bring new skills
 - Distributed data processing and cleaning
 - DevOps and MLOps
- Technologies are improving
 - SQL over large files
 - Self describing file formats & catalog managers
- Organizations are evolving
 - Tracking data usage and file permissions
 - New job title: data engineers



Where to go from here

- Data Warehouse
- Parallel Databases
- Big Data Processing
- Cloud Databases
- Data Science

Why compute in parallel?

- Multi-cores:
 - Most processors have multiple cores
 - This trend will likely increase in the future
- Big data: too large to fit in main memory
 - Distributed query processing on 100x-1000x servers
 - Widely available now using cloud services

Parallel DBMSs

- How to evaluate a parallel DBMS?
- How to architect a parallel DBMS?
- How to partition data in a parallel DBMS?

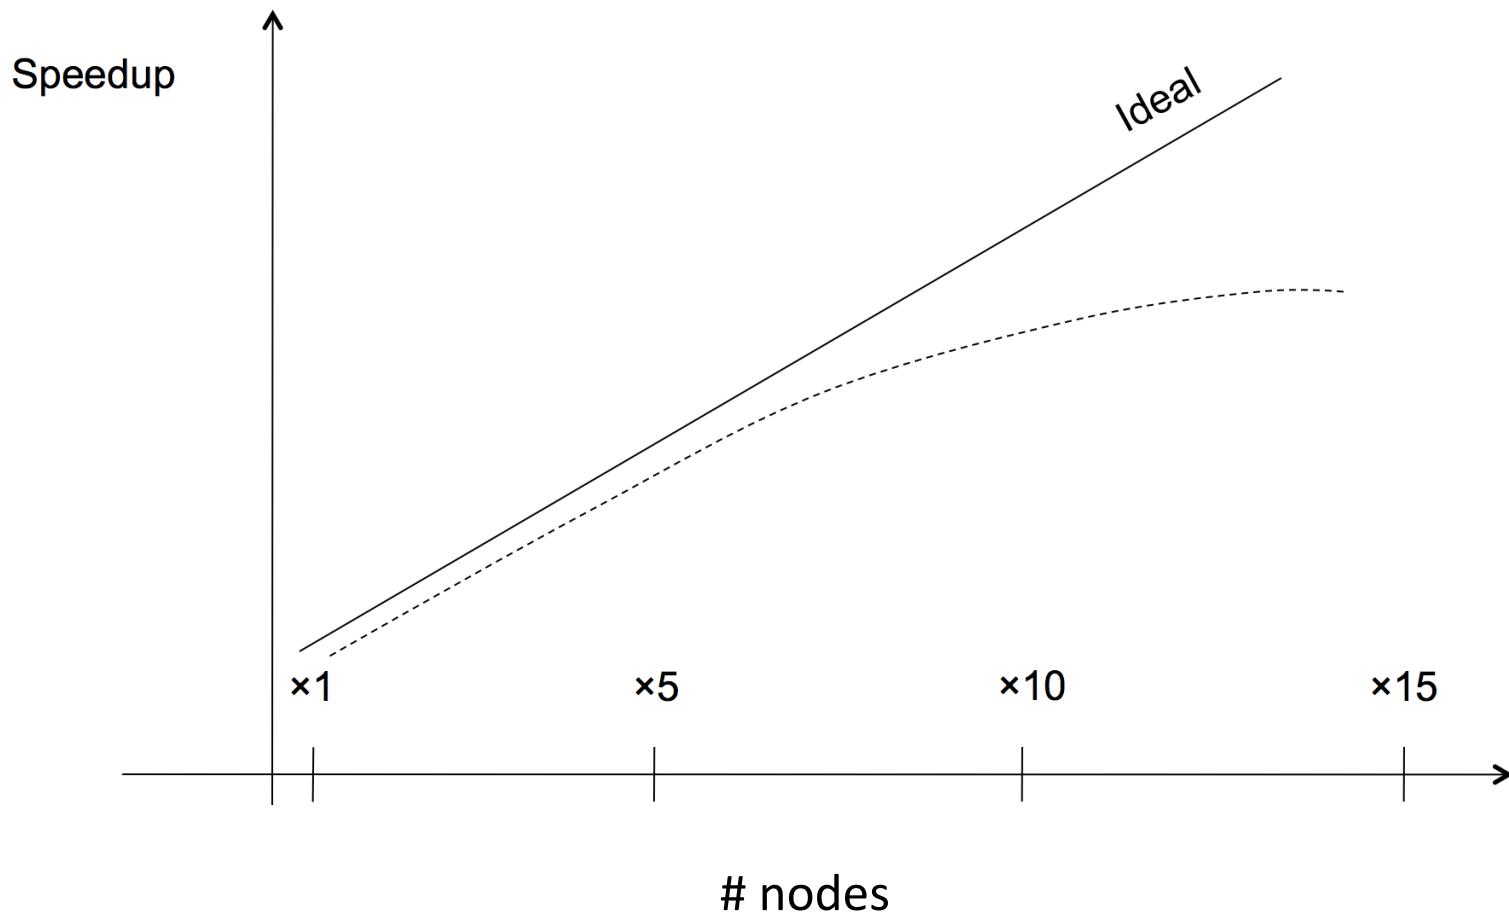
Parallel DBMSs

- **How to evaluate a parallel DBMS?**
- How to architect a parallel DBMS?
- How to partition data in a parallel DBMS?

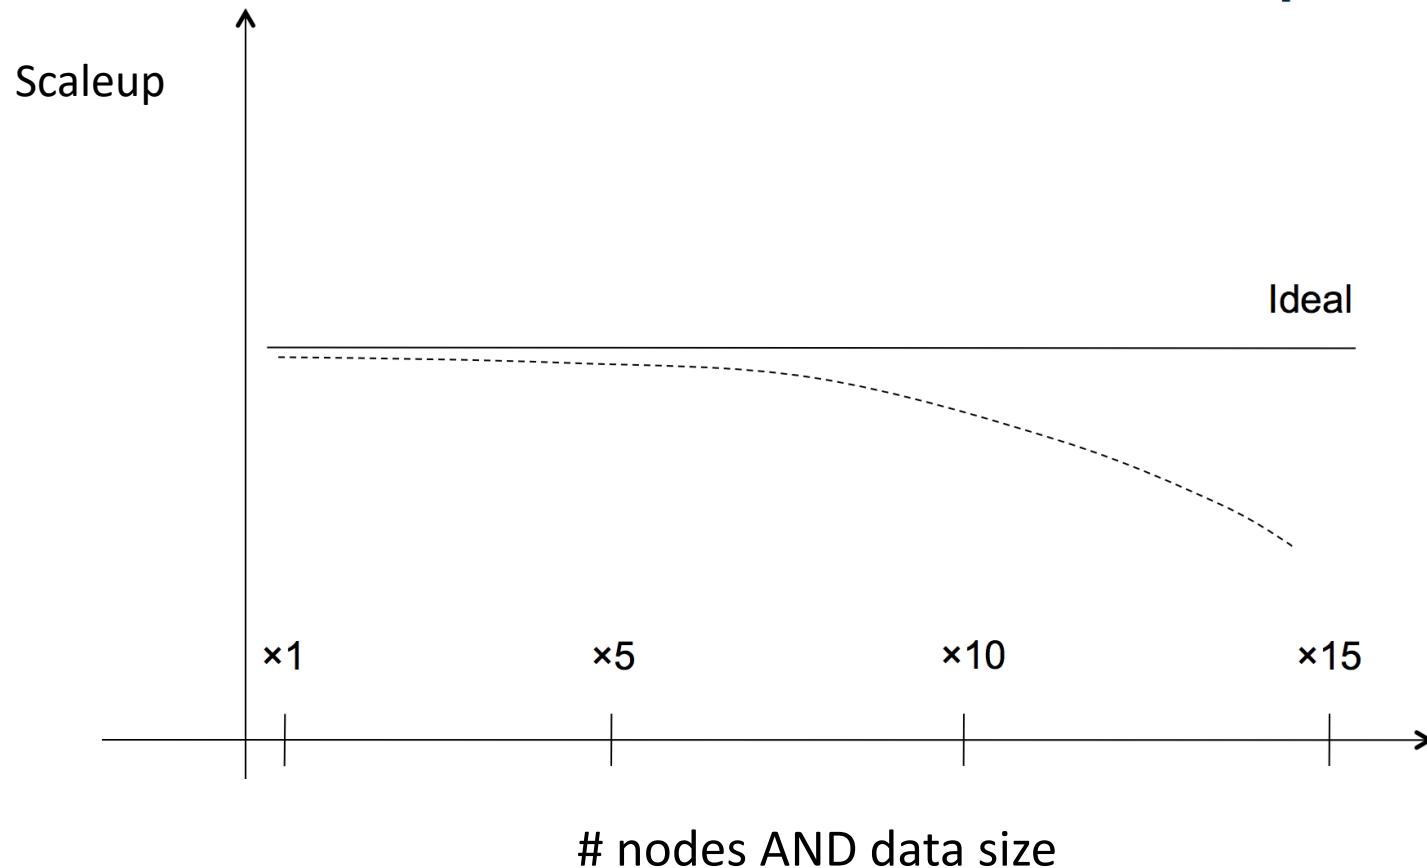
Performance Metrics for Parallel DBMSs

- Nodes = processors, compute
- **Speedup**
 - More nodes, same data → Higher speed
- **Scaleup**
 - More nodes, more data → same speed

Linear v.s. Non-linear Speedup



Linear v.s. Non-linear Scaleup



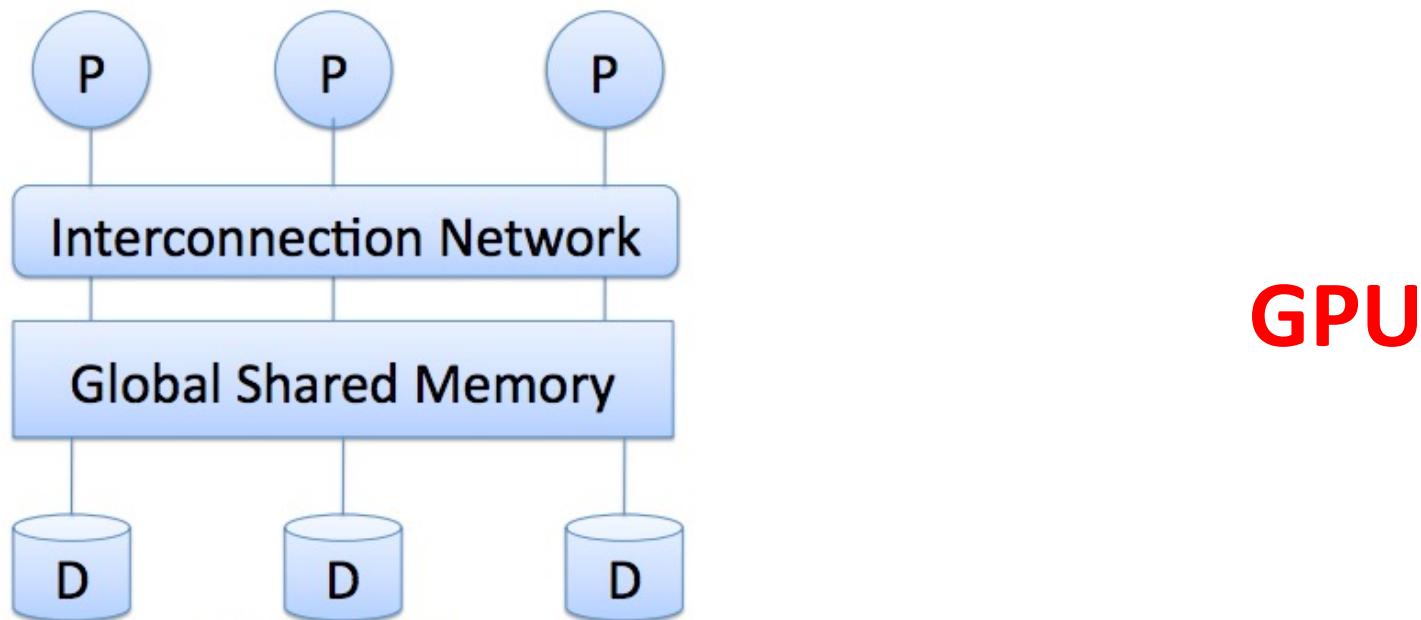
Parallel DBMSs

- How to evaluate a parallel DBMS?
- How to architect a parallel DBMS?
- How to partition data in a parallel DBMS?

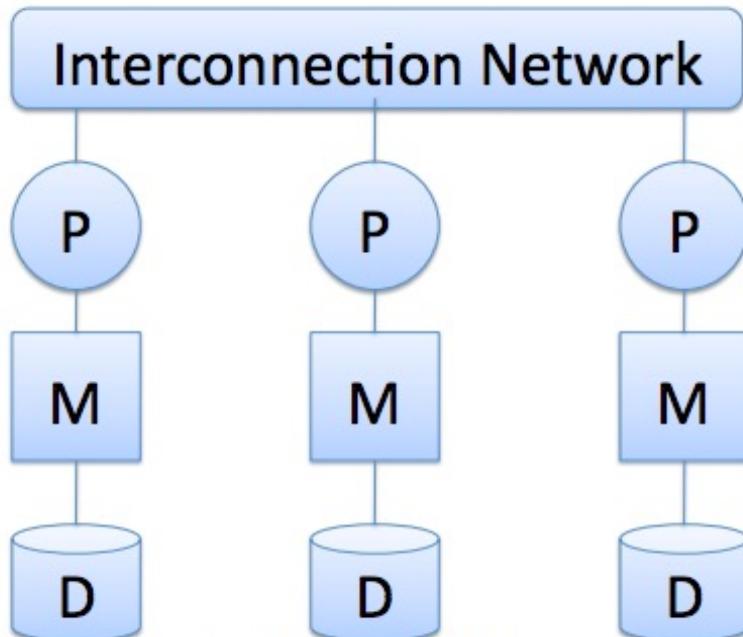
Three Architectures

- **Shared Memory**
- **Shared Nothing**
- **Shared Disk**

Shared Memory

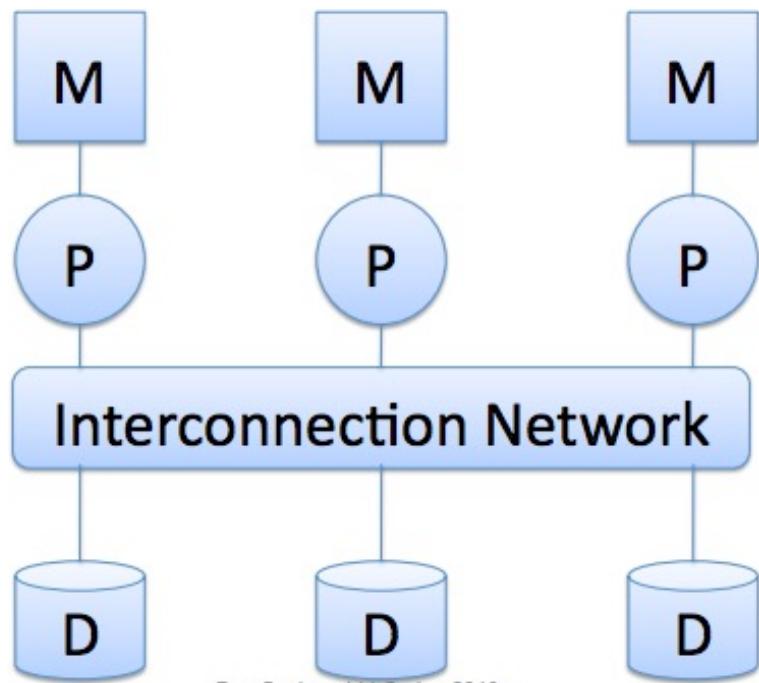


Shared Nothing



Parallel DBMSs, MapReduce,
Spark

Shared Disk



Azure Synapse Analytics

Three Architectures

- **Shared Memory**



Computation vs. Communication Trade-offs

- **Shared Nothing**



Economic Consideration

- **Shared Disk**

Parallel DBMSs

- How to evaluate a parallel DBMS?
- How to architect a parallel DBMS?
- How to partition data in a parallel DBMS?

Horizontal Data Partitioning

- **Round Robin**

- ☺ Load Balancing
- ☹ Bad Query Performance

- **Range Partitioning**

- ☺ Good for range/point queries
- ☹ Data Skew (i.e., Bad Load balancing)

- **Hash Partitioning**

- ☺ Load Balancing, Good for point queries
- ☹ Hard to answer range queries

Where to go from here

- Data Warehouse
- Parallel Databases
- **Big Data Processing**
- Cloud Databases
- Data Science

3 Vs of Big Data

- Volume: data size
- Velocity: rate of data coming in
- Variety: data sources, formats, workloads

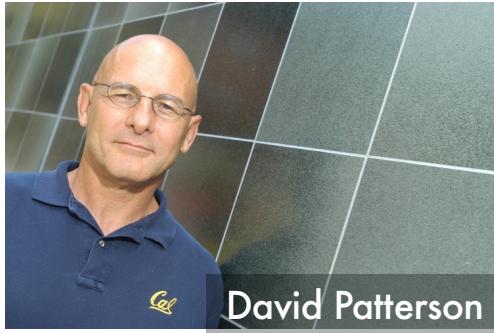
Big Data Systems



Shared Nothing Architecture

How was Spark created?

UC Berkeley's Research Centers



David Patterson

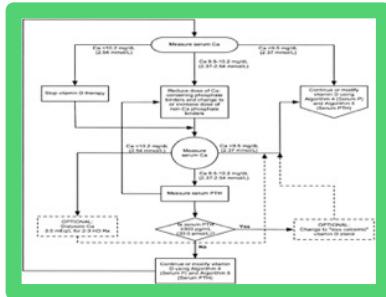
Requirements

- A common vision
- About 5 years
- At least three faculty
- A dozen students

| Years | Title | Profs: Director, Co-PIs | Students |
|-----------|--|--|----------|
| 1977–1981 | X-Tree: Tree Multiprocessor | Despain, Patterson, Sequin | 12 |
| 1980–1984 | RISC: Reduced Instructions | Patterson, Ousterhout, Sequin | 17 |
| 1983–1986 | SOAR: Smalltalk On A RISC | Patterson, Ousterhout | 22 |
| 1985–1989 | SPUR: Symbolic Processing Using RISCs | Patterson, Fateman, Hiltzinger, Hodges, Katz, Ousterhout | 21 |
| 1988–1992 | RAID: Redundant Array of Inexpensive Disks | Katz, Ousterhout, Patterson, Stonebraker | 16 |
| 1993–1998 | NOW: Network of Workstations | Culler, Anderson, Brewer, Patterson | 25 |
| 1997–2002 | IRAM: Intelligent RAM | Patterson, Kubiatowicz, Wawrzynek, Yelick | 12 |
| 2001–2005 | ROC: Recovery Oriented Computing Systems | Patterson, Fox | 11 |
| 2005–2011 | RAD Lab: Reliable Adaptive Distributed Computing Lab | Patterson, Fox, Jordan, Joseph, Katz, Shenker, Stoica | 30 |
| 2007–2013 | Par Lab: Parallel Computing Lab | Patterson, Asanovic, Demmel, Fox, Keutzer, Kubiatowicz, Sen, Yelick | 40 |
| 2011–2017 | AMP Lab: Algorithms, Machines, and People | Franklin, Jordan, Joseph, Katz, Patterson, Recht, Shenker, Stoica | 40 |
| 2013–2018 | ASPIRE Lab | Asanovic, Alon, Bachrach, Demmel, Fox, Keutzer, Nikolic, Patterson, Sen, Wawrzynek | 40 |

AMPLab's Vision

Make sense of BIG DATA by tightly integrating
algorithms, machines, and people



+

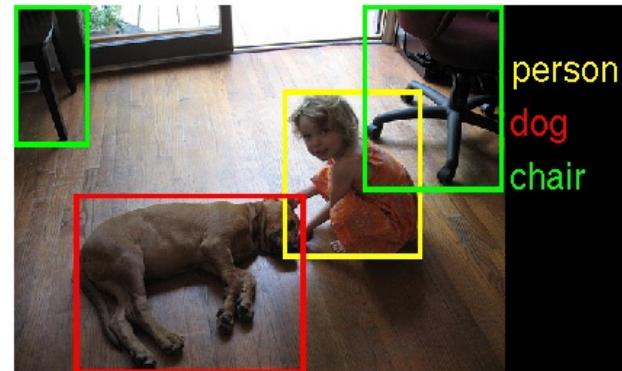


+



Example: Extract Value From Image Data

What are in the image?



How to solve the problem?

Deep Learning (Algorithms)
GPU Cluster (Machines)
ImageNet (People)

Spark's Initial Idea

- Algorithms + Machines
 - Run ML Algorithms on Hadoop
- Why is it slow?
 1. The algorithms are iterative (i.e., multiple scans of data)
 2. MapReduce writes/reads data to/from **disk** at each iteration
- Solution
 - Keep data in **memory**



Matei Zaharia

How About Fault Tolerance?

Resilient Distributed Datasets (RDD)



Main Idea: Logging the transformations (used to build an RDD) rather than the RDD itself

Why Spark?

Fast

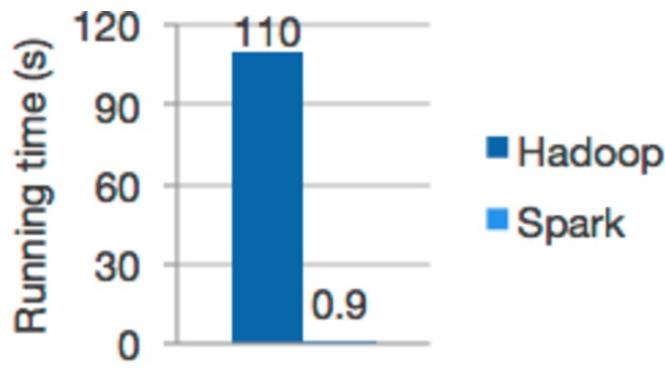


Easy to Use



What Makes Spark *Fast*?

- In-memory Computation



Logistic regression in Hadoop and Spark

What you save?

- Serialization/Deserialization
- Compression/Decompression
- I/O cost



“CPU (and not I/O) is often the bottleneck”

Ousterhout et al. Making Sense of Performance in Data Analytics Frameworks. NSDI 2015: 293-307

Hardware Trends

| | 2010 | 2016 | |
|---------|------------------|-------------------|-----|
| Storage | 50+MB/s (HDD) | 500+MB/s (SSD) | 10X |
| Network | 1Gbps | 10Gbps | 10X |
| CPU | ~3GHz | ~3GHz | |

What Makes Spark *Fast*?

Project Tungsten: Bringing Apache Spark Closer to Bare Metal



by Reynold Xin and Josh Rosen

Posted in **ENGINEERING BLOG** | April 28, 2015

1. Memory Management and Binary Processing
2. Cache-aware computation
3. Code generation

Why Spark?

Fast



Easy to Use



What Makes Spark Easy-to-Use?

- Over 80 High-level Operators

WordCount (Mapreduce)

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;

public class WordCount extends Configured implements Tool {

    public static class TokenizerMapper
        extends Mapper<LongWritable, Text, Text, IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        @Override
        public void map(LongWritable key, Text value, Context context
            ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }

    public static class IntSumReducer
        extends Reducer<Text, IntWritable, Text, IntWritable> {
        private IntWritable result = new IntWritable();

        @Override
        public void reduce(Text key, Iterable<IntWritable> values,
                           Context context
                           ) throws IOException, InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            result.set(sum);
            context.write(key, result);
        }

        public static void main(String[] args) throws Exception {
            int res = ToolRunner.run(new Configuration(), new WordCount(), args);
            System.exit(res);
        }

        @Override
        public int run(String[] args) throws Exception {
            Configuration conf = this.getConf();
            Job job = Job.getInstance(conf, "word count");
            job.setJarByClass(WordCount.class);
            job.setInputFormatClass(TextInputFormat.class);
            job.setMapperClass(TokenizerMapper.class);
            job.setCombinerClass(IntSumReducer.class);
            job.setReducerClass(IntSumReducer.class);

            job.setOutputKeyClass(Text.class);
            job.setOutputValueClass(IntWritable.class);
            job.setOutputFormatClass(TextOutputFormat.class);
            TextInputFormat.addInputPath(job, new Path(args[0]));
            TextOutputFormat.setOutputPath(job, new Path(args[1]));

            return job.waitForCompletion(true) ? 0 : 1;
        }
    }
}
```

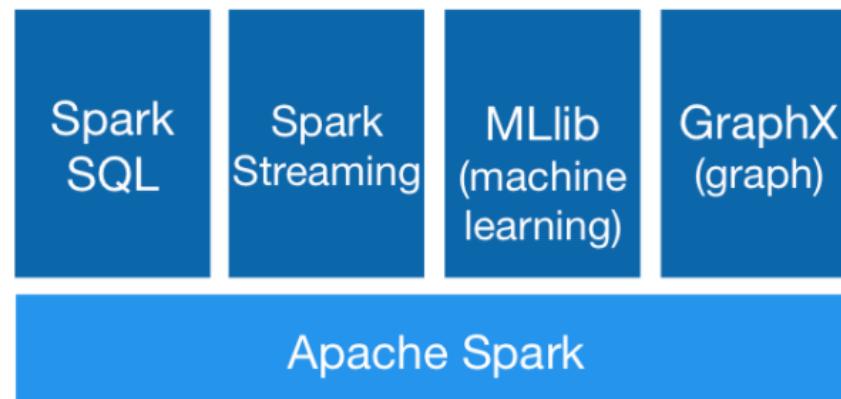
WordCount (Spark)

```
text_file = spark.textFile("hdfs://...")

text_file.flatMap(lambda line: line.split())
    .map(lambda word: (word, 1))
    .reduceByKey(lambda a, b: a+b)
```

What Makes Spark *Easy-to-Use*?

- Unified Engine



Easy to manage, learn, and combine functionality

Analogy



Specialized Devices



Unified Device

What Makes Spark *Easy-to-Use*?

- Integrate Broadly

Languages:



The Big Data world is diversified



The Apache Spark logo is displayed prominently in the center. It features the word "Spark" in a large, dark gray, sans-serif font. Above the letter "S", there is a smaller, semi-transparent "APACHE" logo. A five-pointed orange star is positioned above the letter "k". Below the "Spark" text, there is a small trademark symbol (TM).

Data Sources:



Environments:



Where to go from here

- Data Warehouse
- Parallel Databases
- Big Data Processing
- Cloud Databases
- Data Science

Amazon

- From Wikipedia 2006

Article | Talk | Head

Amazon.com

From Wikipedia, the free encyclopedia

This is an old revision of this page
05:10, 20 March 2006 ([→Customer permanent link](#) to this revision, wh revision.

(diff) ← Previous revision | Latest revisi

Amazon.com

(NASDAQ: AMZN) is an American electronic commerce company based in Seattle, Washington. It was one of the

- From Wikipedia 2016

Amazon.com

From Wikipedia, the free encyclopedia

Further information: [Timeline of Amazon](#)

Amazon.com, Inc. (/əməzən/ or /əməzən/), often referred to as simply **Amazon**, is an American electronic commerce and cloud computing company with headquarters in Seattle, Washington. It is the

What is Cloud Computing?

- The buzz word before “Big Data”
 - Larry Ellison’s response in 2009 (<https://youtu.be/UOEFXaWHppE?t=7s>)
 - Berkeley RADLab’s paper in 2009 (<https://youtu.be/IJCxqoh5ep4>)
- A technical point of view
 - Internet-based computing (i.e., computers attached to network)
- A business-model point of view
 - Pay-as-you-go (i.e., rental)



Three Types of Cloud Computing

CourSys

Application + Cloud = **SaaS** (Software as a service)

Database

Platform + Cloud = **PaaS** (Platform as a service)

Servers

Infrastructure + Cloud = **IaaS** (Infrastructure as a service)

Cloud Databases

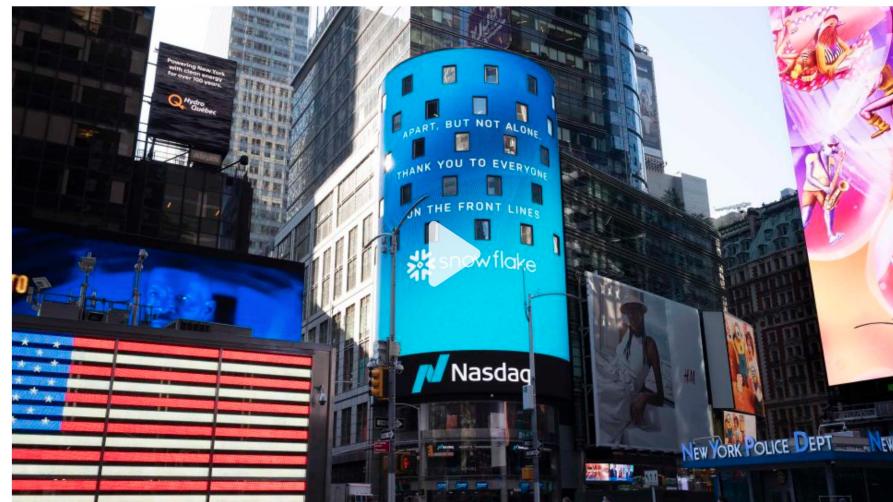
- Amazon
 - Aurora
 - Redshift
- Microsoft
 - Azure SQL Databases
 - Azure Synapse Analytics
- Google
 - Spanner
 - BigQuery

Snowflake shares more than double. It's the biggest software IPO ever



By [Paul R. La Monica](#), CNN Business

Updated 3:37 AM EDT, Thu September 17, 2020



Where to go from here

- Data Warehouse
- Parallel Databases
- Big Data Processing
- Cloud Databases
- **Data Science**

Summary

- Data Warehouse
- Parallel Databases
- Big Data Processing
- Cloud Databases

Acknowledge

- Some lecture slides were copied from or inspired by the following course materials
 - “W4111: Introduction to databases” by Eugene Wu at Columbia University
 - “CSE344: Introduction to Data Management” by Dan Suciu at University of Washington
 - “CMPT354: Database System I” by John Edgar at Simon Fraser University
 - “CS186: Introduction to Database Systems” by Joe Hellerstein at UC Berkeley
 - “CS145: Introduction to Databases” by Peter Bailis at Stanford
 - “CS 348: Introduction to Database Management” by Grant Weddell at University of Waterloo
 - UC Berkeley DS100 Fall 2017